

DELE CA2

Tan Yu Hoe

P2026309

Diploma in Applied AI and Analytics

DAAA/FT/2B/04

School of Computing

Singapore Polytechnic

tanyh.20@ichat.sp.edu.sg

DELE

ST1504

CA2 Assignment



Table of Contents

- **Methodology and Motivation**
- **Exploratory Data Analysis**
- **Feature Engineering and Background Research**
- **Deep Convolutional GAN**
- **Conditional GAN**
- **Auxiliary Classifier GAN**
- **Conclusion**



Methodology and Motivation

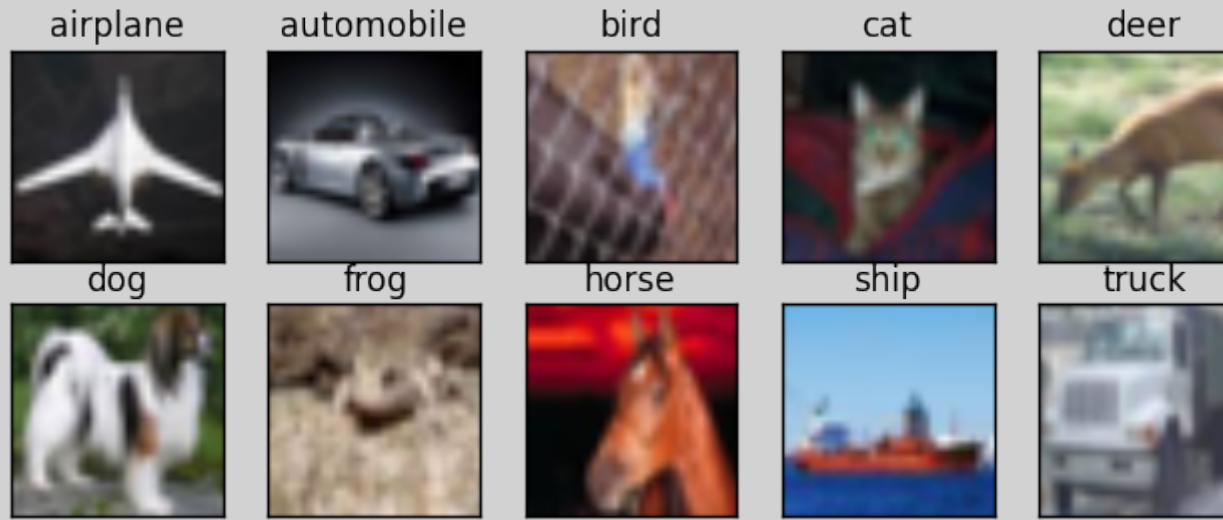
In this assignment, Part A requires us to build and train a GAN architecture to create **1000 small colour images**, based on the CIFAR 10 dataset.

To attempt this task, I will be exploring various GAN architectures such as CGAN and ACGAN produce synthetic images of the CIFAR 10 dataset of substantiable quality.





CIFAR-10

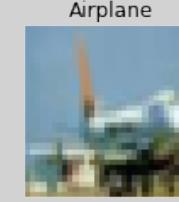
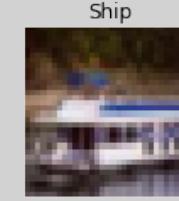
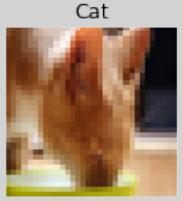
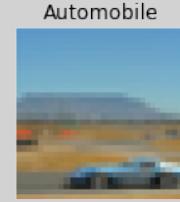
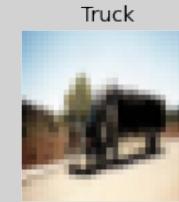


CIFAR-10 dataset (Canadian Institute for Advanced Research) is a collection of images that are commonly used to train and benchmark image classification algorithms. It is a subset of 80 million tiny images and consists of 60,000 instances - 32 by 32 coloured images from 10 different classes. There are 50,000 training examples and 10,000 examples.



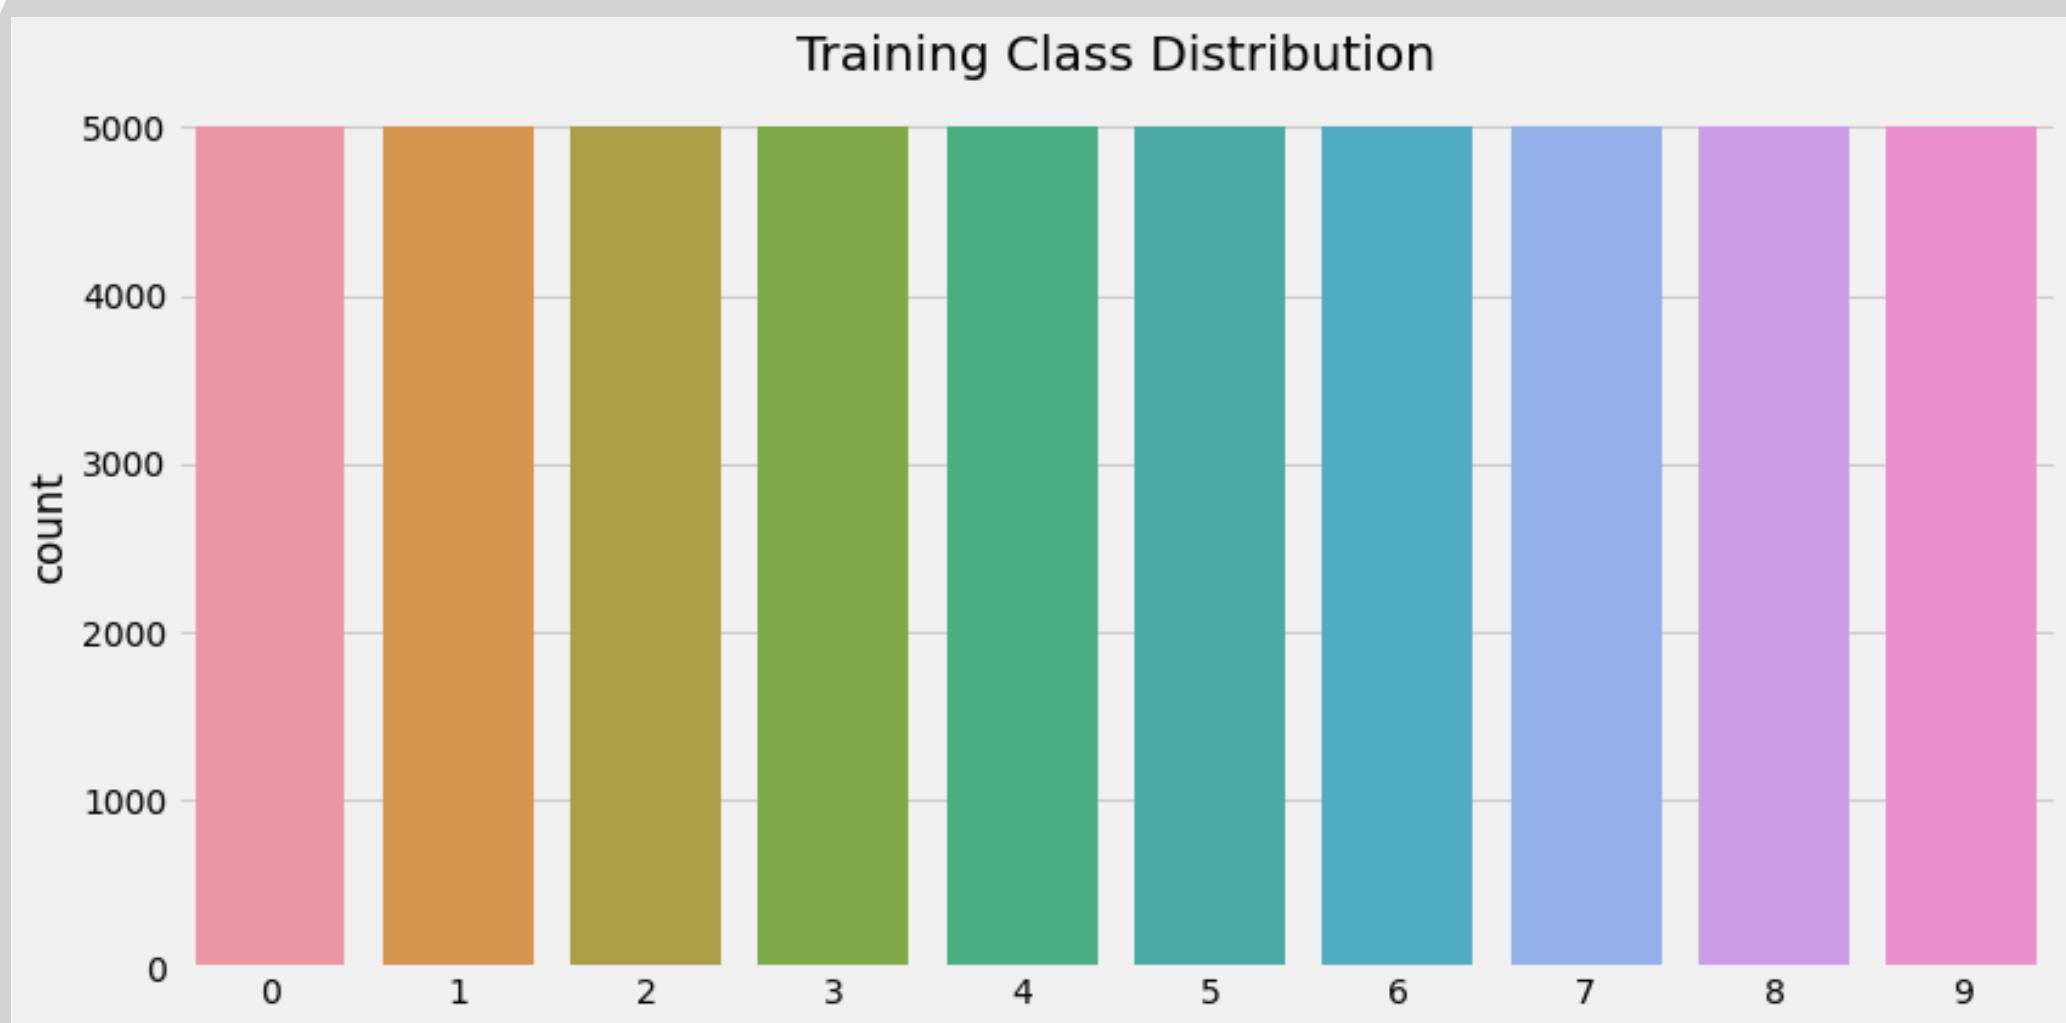


CIFAR-10





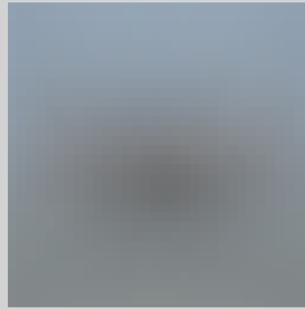
CIFAR-10





CIFAR-10

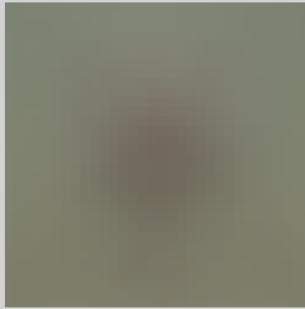
μ Airplane



μ Automobile



μ Bird



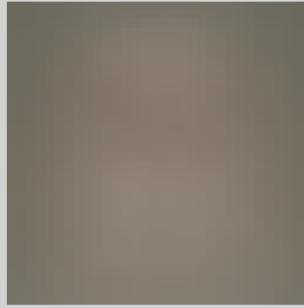
μ Cat



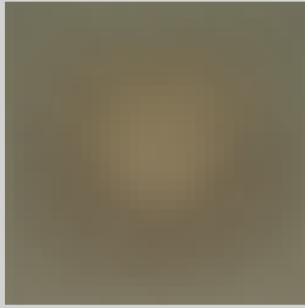
μ Deer



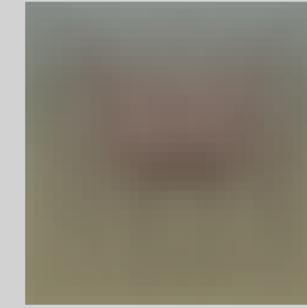
μ Dog



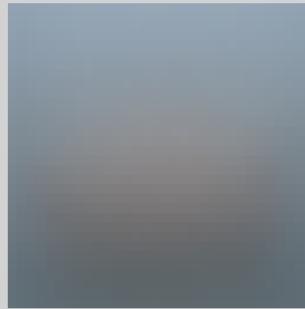
μ Frog



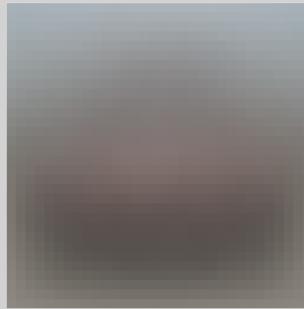
μ Horse



μ Ship



μ Truck



Generative Adversarial Networks on CIFAR-10

CIFAR-10 is a dataset of classes, with images and labels provided. I went to research for papers of the following properties:

- Generative Adversarial Networks utilising Labels
- Improvements on reliability on Generative Adversarial Networks training



Generative Adversarial Networks on CIFAR-10

For the first property, I found a 2014 paper on “*Conditional Adversarial Net*”. [1] The paper discussed about using label y as “conditions”, the term coming from conditional probability $P(A|B)$, to be fed into the generator and discriminator.

The results from the paper have shown that using conditions for image generation have greatly aid in image generation for MNIST digits dataset, outperforming non-conditional adversarial networks during that era.

1. Mirza, M., & Osindero, S. (2014). Conditional Generative Adversarial Nets. arXiv [cs.LG]. Opgehaal van <http://arxiv.org/abs/1411.1784>



Generative Adversarial Networks on CIFAR-10

For the second property, I found a paper from 2016 on “*Conditional Image Synthesis with Auxiliary Classifier GANs*”. [1] The paper describes about extending a CGAN architecture to an AC-GAN architecture, by adding label conditioning.

The author explained that adding label conditioning improved image discrimination by ignoring the component of loss arising from class labels when there are no labels for training unsupervised GAN.

1. Odena, A., Olah, C., & Shlens, J. (2017). Conditional Image Synthesis With Auxiliary Classifier GANs. arXiv [stat.ML]. Opgehaal van <http://arxiv.org/abs/1610.09585>





Pixel Normalization / Rescaling

The purpose of normalisation in image processing is to “attempt” to bring the pixels down to a normal distribution $N(0,1)$, to mitigate the strong influence of very large or very small pixels. Here, I attempt to use **Min-Max Normalization** (also called Unity-Based Normalization), to bring the pixel values down to range $[0, -1]$. This is a necessary pre-processing step to improve better optimization within the neural network.

$$X' = a + \frac{(X - X_{min})(b - a)}{X_{max} - X_{min}} = -1 + \frac{X}{127.5}$$

where a and b are lower and upper bound of a predefined range.

For GAN, it is recommended to use hyperbolic tangent **tanh** activation function in the output layer from the generator model. Given the saturating activation function provides an output range of $[-1, 1]$, it would be recommended to scale down the image to a range $[-1, 1]$.





Deep Convolutional GAN

I started with a **Deep Convolutional GAN** (DCGAN) as a Baseline Model. The architecture consists of two neural networks – (1) a **generator** and (2) a **discriminator**.

1. **Generator G** : to generate the synthetic images $G(z)$ from a latent noise vector (z) [as a Gaussian Noise $N(0,1)$]
2. **Discriminator D** : to discriminate synthetic images $G(z)$ from the training images x



Deep Convolutional GAN Discriminator

The discriminator is used to “classify” (the better term would be “discriminate”) between a Real Image and a Synthetic Image. The network structure follows a typical CNN architecture. I have made some changes to the network; however some changes were found to be beneficial while some are found to be disadvantageous.

- 2 by 2 Strided Convolutions
- Leaky ReLU Activation Function - $\alpha = 0.2$
- Batch Normalization
- Spectral Normalization
- Gaussian Weights Initialization - $N(0, 0.2)$



Model: "discriminator_GAN"

Layer (type)	Output Shape	Param #
<hr/>		
spectral_normalization (SpectralNormalization)	(None, 16, 16, 64)	3200
batch_normalization (BatchNormalization)	(None, 16, 16, 64)	256
leaky_re_lu (LeakyReLU)	(None, 16, 16, 64)	0
spectral_normalization_1 (SpectralNormalization)	(None, 8, 8, 128)	131328
batch_normalization_1 (BatchNormalization)	(None, 8, 8, 128)	512
leaky_re_lu_1 (LeakyReLU)	(None, 8, 8, 128)	0
spectral_normalization_2 (SpectralNormalization)	(None, 4, 4, 256)	524800
batch_normalization_2 (BatchNormalization)	(None, 4, 4, 256)	1024
leaky_re_lu_2 (LeakyReLU)	(None, 4, 4, 256)	0
spectral_normalization_3 (SpectralNormalization)	(None, 2, 2, 512)	2098176
batch_normalization_3 (BatchNormalization)	(None, 2, 2, 512)	2048
leaky_re_lu_3 (LeakyReLU)	(None, 2, 2, 512)	0
global_max_pooling2d (GlobalMaxPooling2D)	(None, 512)	0
dense (Dense)	(None, 1)	513
<hr/>		
Total params:	2,761,857	
Trainable params:	2,758,977	
Non-trainable params:	2,880	

```
# function to create discriminator model
def create_discriminator(image_size):
    weights_init = RandomNormal(mean=0, stddev=0.02)
    discriminator = Sequential([
        Input(shape=input_dim),
        ## 2 by 2 Strides for downsampling
        ## Leaky ReLU for Activation
        SpectralNormalization(
            Conv2D(64, kernel_size=4, strides=2, padding='same', kernel_initializer=weights_init)
        ),
        BatchNormalization(momentum=0.8),
        LeakyReLU(0.2),
        # Conv 2
        SpectralNormalization(
            Conv2D(128, kernel_size=4, strides=2, padding='same', kernel_initializer=weights_init)
        ),
        BatchNormalization(momentum=0.8),
        LeakyReLU(0.2),
        # Conv 3
        SpectralNormalization(
            Conv2D(256, kernel_size=4, strides=2, padding='same', kernel_initializer=weights_init)
        ),
        BatchNormalization(momentum=0.8),
        LeakyReLU(0.2),
        # # Conv 4
        SpectralNormalization(
            Conv2D(512, kernel_size=4, strides=2, padding='same', kernel_initializer=weights_init)
        ),
        BatchNormalization(momentum=0.8),
        LeakyReLU(0.2),
        GlobalMaxPooling2D(),
        Dense(1, activation='sigmoid'),
    ], name='discriminator_GAN')
    return discriminator
create_discriminator(image_size=input_dim).summary()
```



Deep Convolutional GAN Generator

The Generator can be thought of as a mapping function from a latent noise vector to a generated image/3D tensor. Some key features I have implemented:

- Transposed Convolution Layer
- Batch Normalization
- Gaussian weights Initialization
- PReLU Activation Function



Model: "generator_GAN"

Layer (type)	Output Shape	Param #
<hr/>		
dense_1 (Dense)	(None, 2048)	264192
reshape (Reshape)	(None, 2, 2, 512)	0
conv2d_transpose (Conv2DTranspose)	(None, 4, 4, 256)	2097408
batch_normalization_4 (BatchNormalization)	(None, 4, 4, 256)	1024
p_re_lu (PReLU)	(None, 4, 4, 256)	4096
conv2d_transpose_1 (Conv2DTranspose)	(None, 8, 8, 128)	524416
batch_normalization_5 (BatchNormalization)	(None, 8, 8, 128)	512
p_re_lu_1 (PReLU)	(None, 8, 8, 128)	8192
conv2d_transpose_2 (Conv2DTranspose)	(None, 16, 16, 64)	131136
batch_normalization_6 (BatchNormalization)	(None, 16, 16, 64)	256
p_re_lu_2 (PReLU)	(None, 16, 16, 64)	16384
conv2d_transpose_3 (Conv2DTranspose)	(None, 32, 32, 3)	3075
<hr/>		
Total params: 3,050,691		
Trainable params: 3,049,795		
Non-trainable params: 896		

```
# function to create generator model
def create_generator(latent_dim):

    # gaussian weights initialization
    weights_init = RandomNormal(mean=0, stddev=0.02)

    generator = Sequential([
        # fc block: handling latent vector z
        Input(shape=(latent_dim,)),
        Dense(2*2*512),
        Reshape((2, 2, 512)),

        # Conv 1
        ## using a kernel size that is a factor of the stride
        Conv2DTranspose(256, kernel_size=4, strides=2, padding='same', kernel_initializer=weights_init),
        BatchNormalization(momentum=0.8),
        PReLU(),

        # Conv 2
        Conv2DTranspose(128, kernel_size=4, strides=2, padding='same', kernel_initializer=weights_init),
        BatchNormalization(momentum=0.8),
        PReLU(),

        # Conv 3
        Conv2DTranspose(64, kernel_size=4, strides=2, padding='same', kernel_initializer=weights_init),
        BatchNormalization(momentum=0.8),
        PReLU(),

        # Output Layer
        Conv2DTranspose(3, kernel_size=4, strides=2, padding='same', activation='tanh', kernel_initializer=weights_init)
    ], name='generator_GAN')
    return generator

create_generator(latent_dim=latent_dim).summary()
```



DCGAN Training Function

```
● ● ●  
  
# Train the discriminator  
with tf.GradientTape() as tape:  
    predictions = self.discriminator(combined_images)  
    d_loss = self.loss_fn(labels, predictions)  
grads = tape.gradient(d_loss, self.discriminator.trainable_weights)  
self.d_optimizer.apply_gradients(  
    zip(grads, self.discriminator.trainable_weights)  
)
```

```
● ● ●  
  
# Train the generator (note that we should *not* update the weights  
# of the discriminator)!  
with tf.GradientTape() as tape:  
    generated_images = self.generator(random_latent_vectors)  
    predictions = self.discriminator(generated_images)  
    g_loss = self.loss_fn(misleading_labels, predictions)  
grads = tape.gradient(g_loss, self.generator.trainable_weights)  
self.g_optimizer.apply_gradients(zip(grads, self.generator.trainable_weights))
```



Deep Convolutional GAN

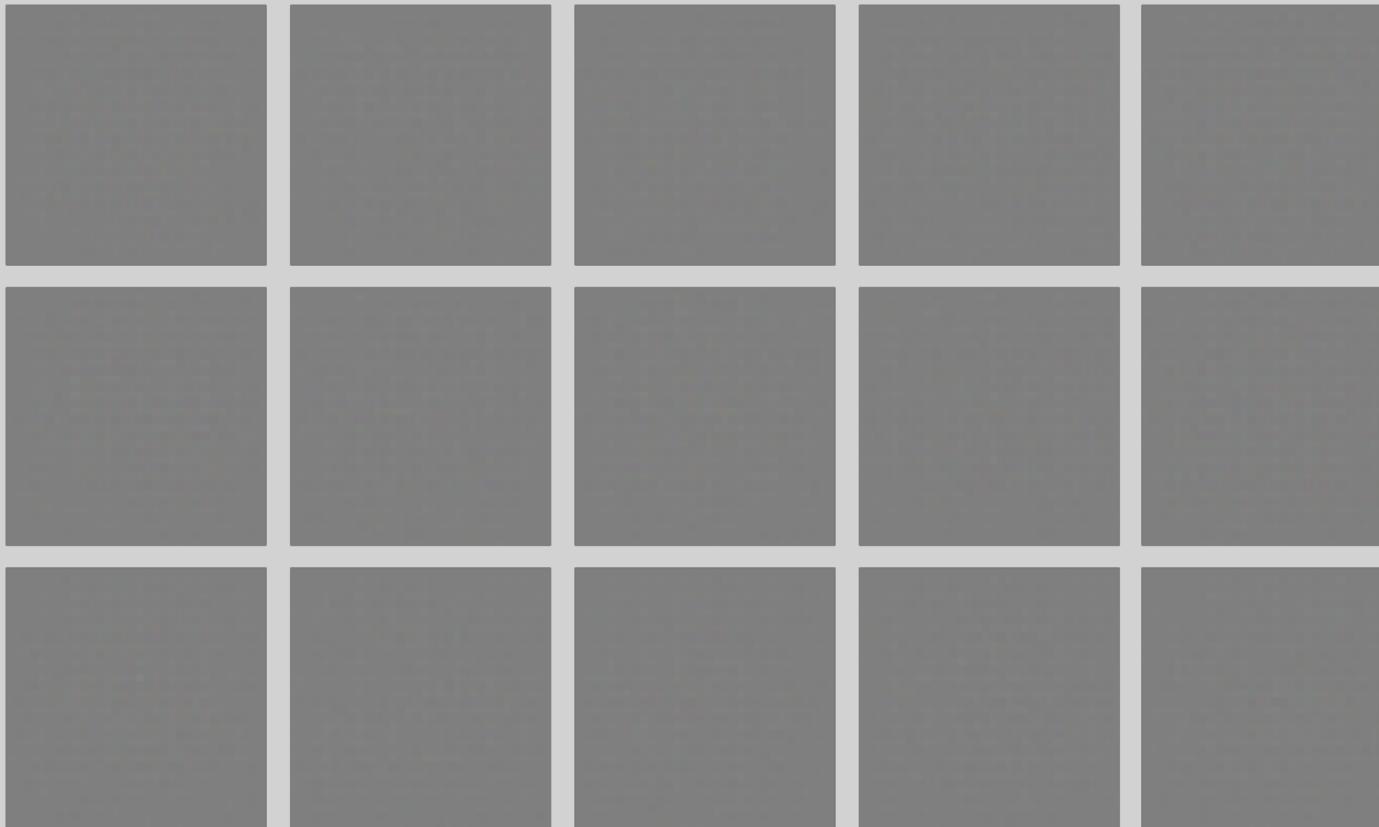
Hyperparameters

- Generator Optimizer : $Adam(\alpha = 0.0002, \beta_1 = 0.5, \beta_2 = 0.999)$
- Discriminator Optimizer : $Adam(\alpha = 0.0002, \beta_1 = 0.5, \beta_2 = 0.999)$
- Batch Size : 64
- Latent Vector Dimensions : 128
- Epochs: 100



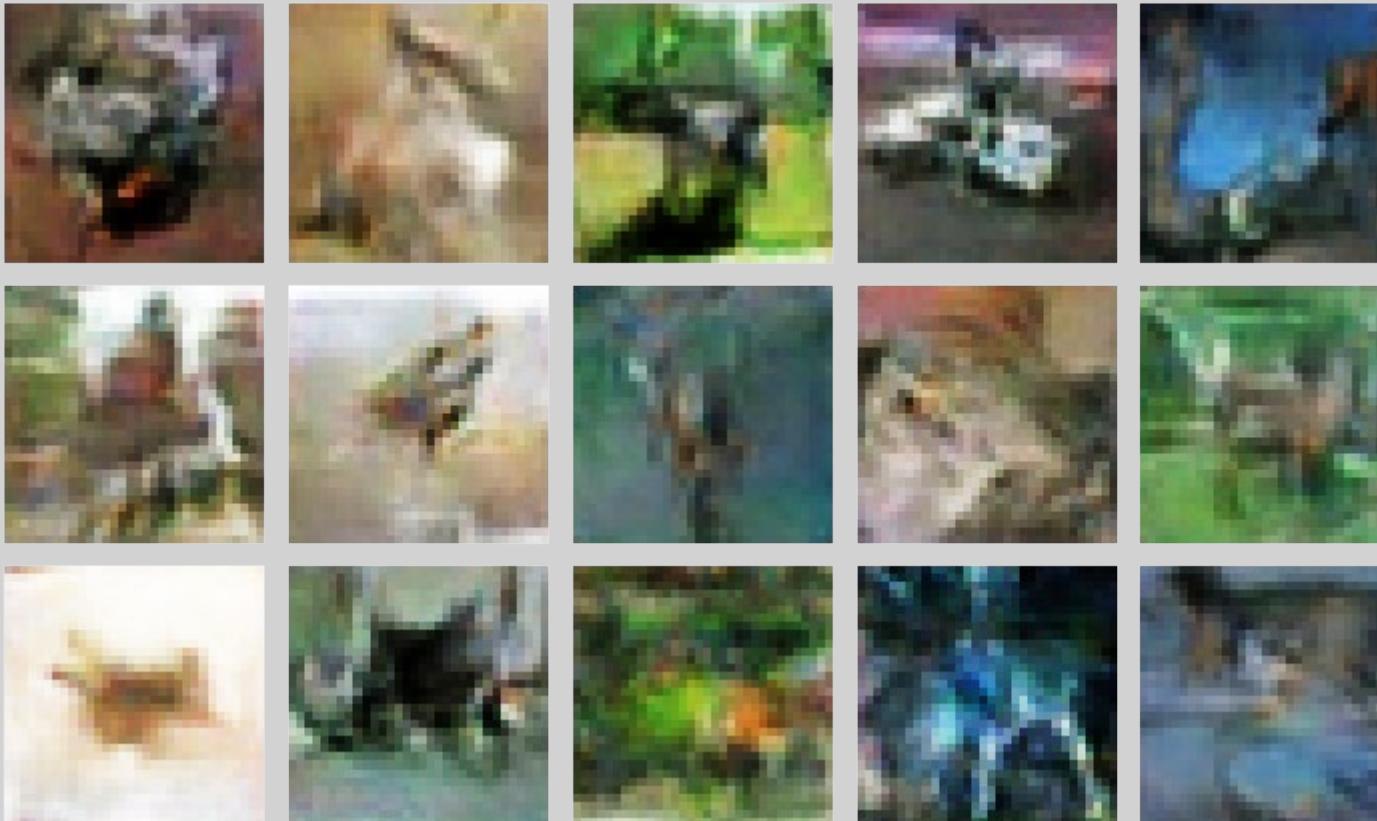
Deep Convolutional GAN

Epoch 0



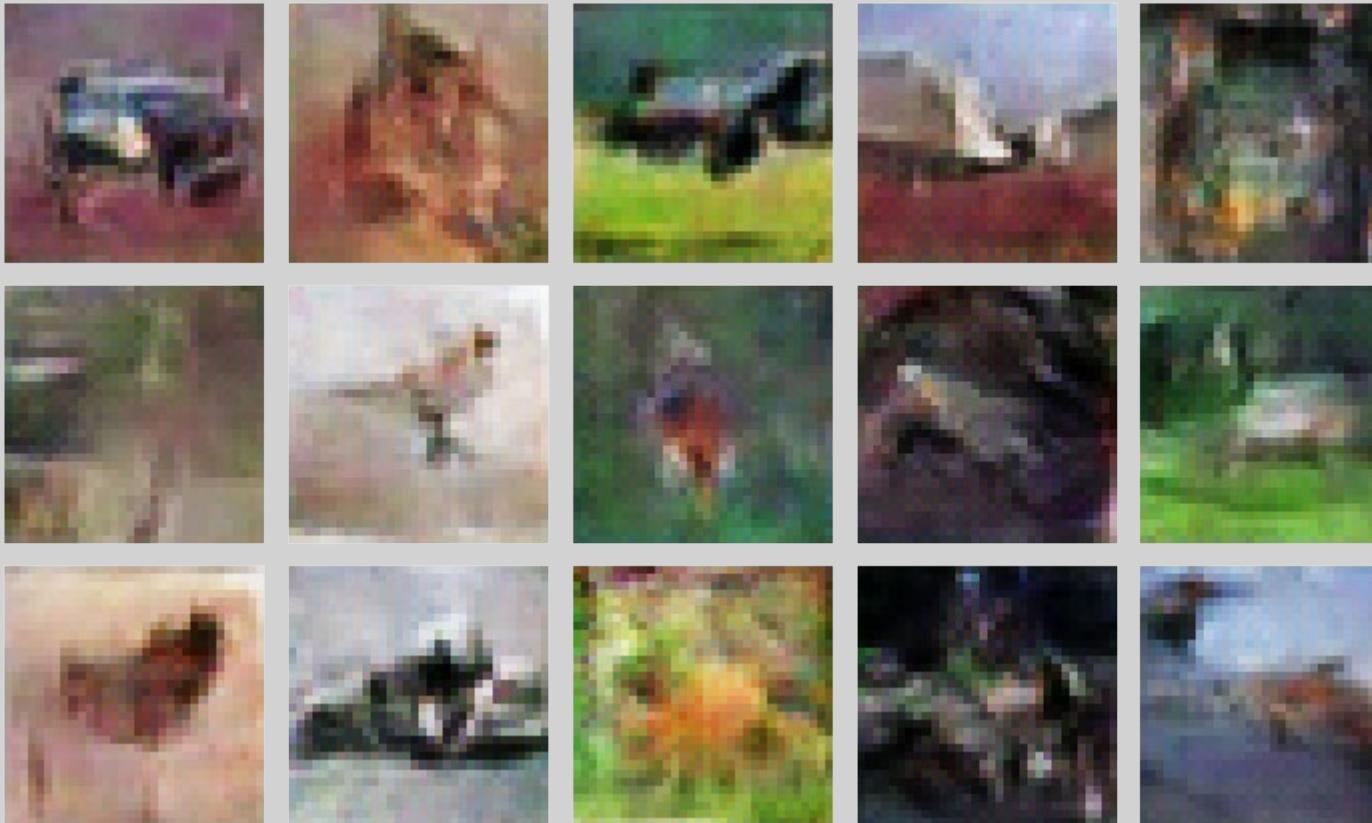
Deep Convolutional GAN

Epoch 25



Deep Convolutional GAN

Epoch 50



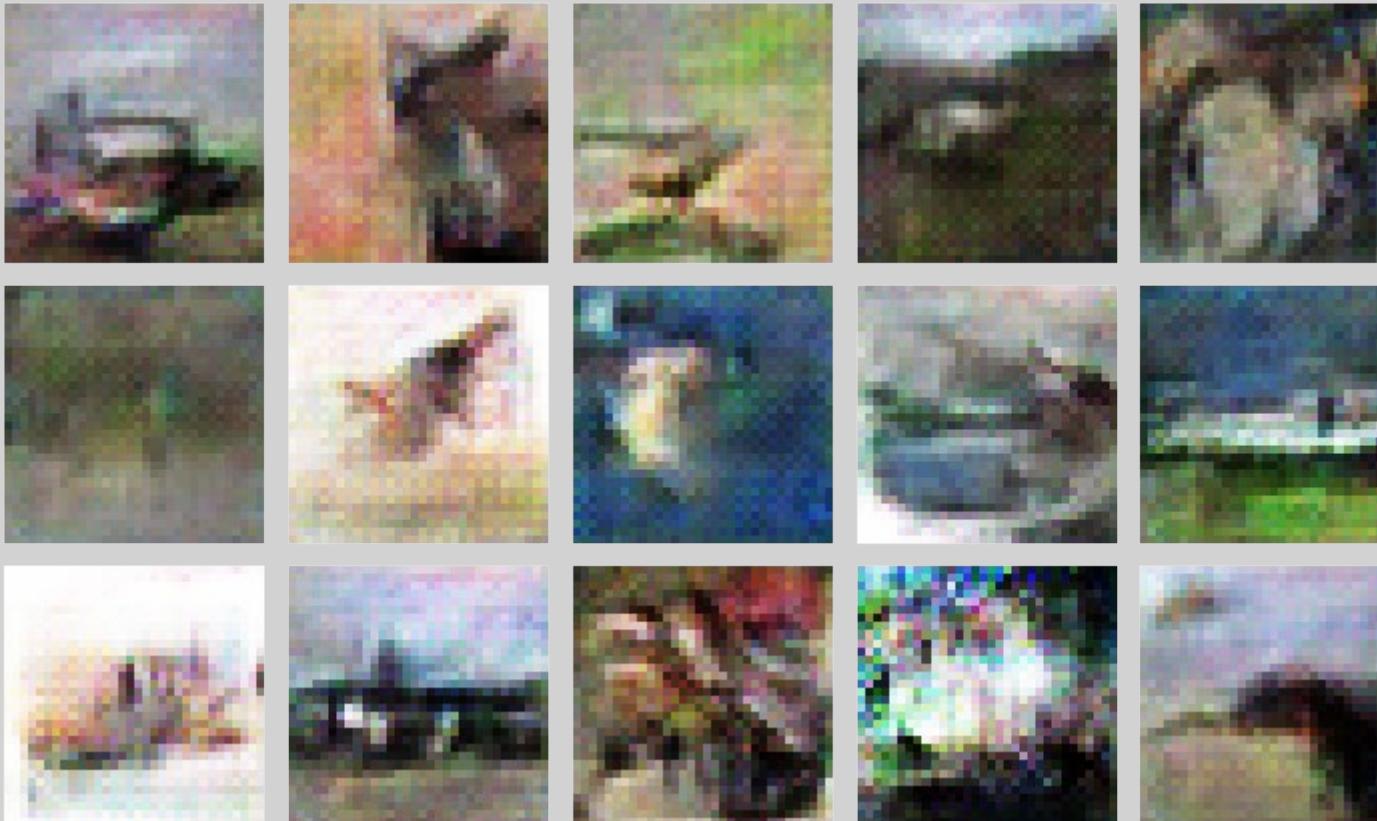
Deep Convolutional GAN

Epoch 75

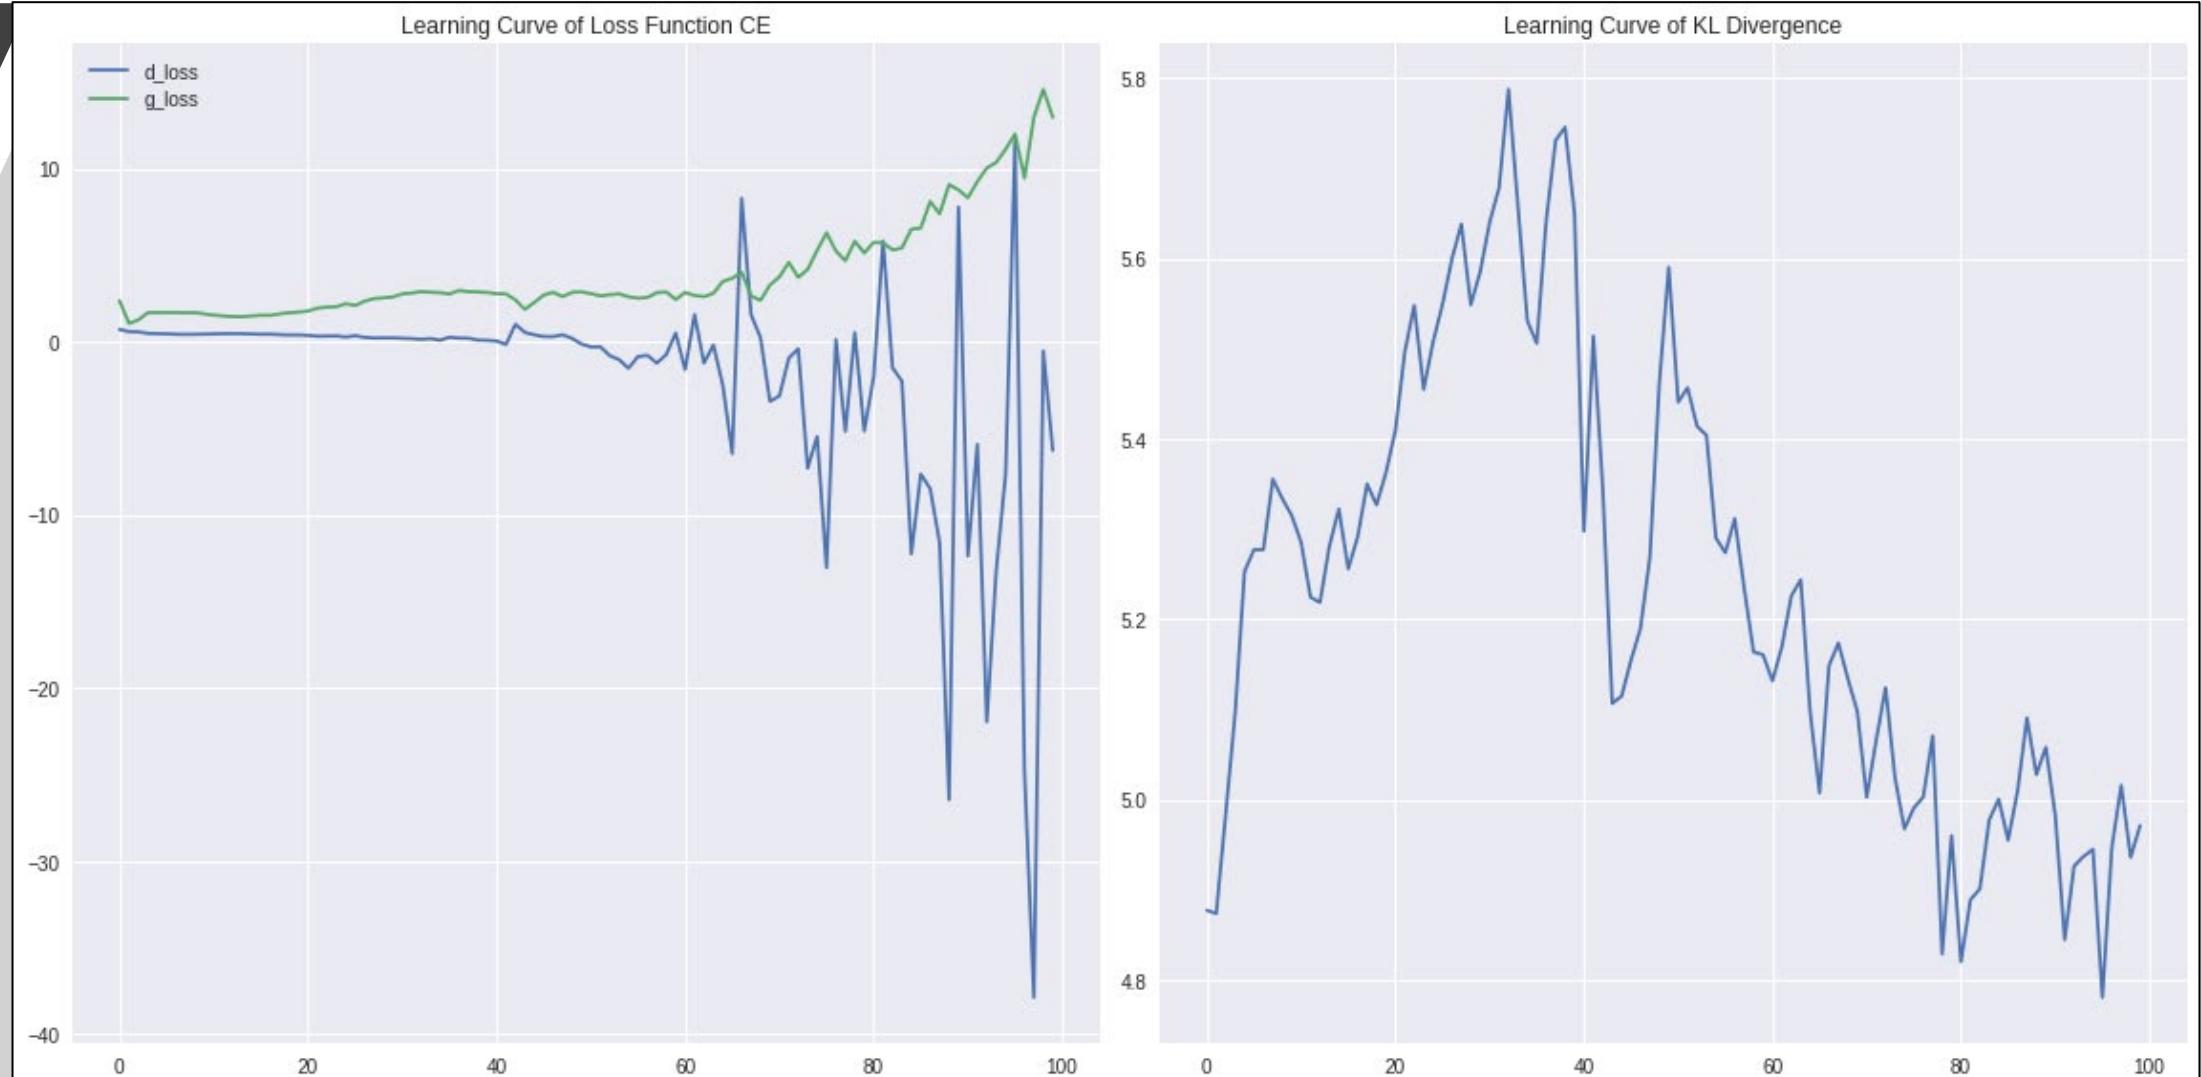


Deep Convolutional GAN

Epoch 100



Deep Convolutional GAN Learning Curve



Deep Convolutional GAN

FID Evaluation

Fréchet Inception Distance is a metric used to access the quality of images created by the generator by comparing the distribution of the generated images with the distribution of real images that were used to train the generator.

To implement FID, the paper on “*GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium*” mentioned that the minimum number of generated sample is 10,000 in order to evaluate FID. I randomly sampled 10,000 images from the training set and generate 10,000 images and evaluate against each other.

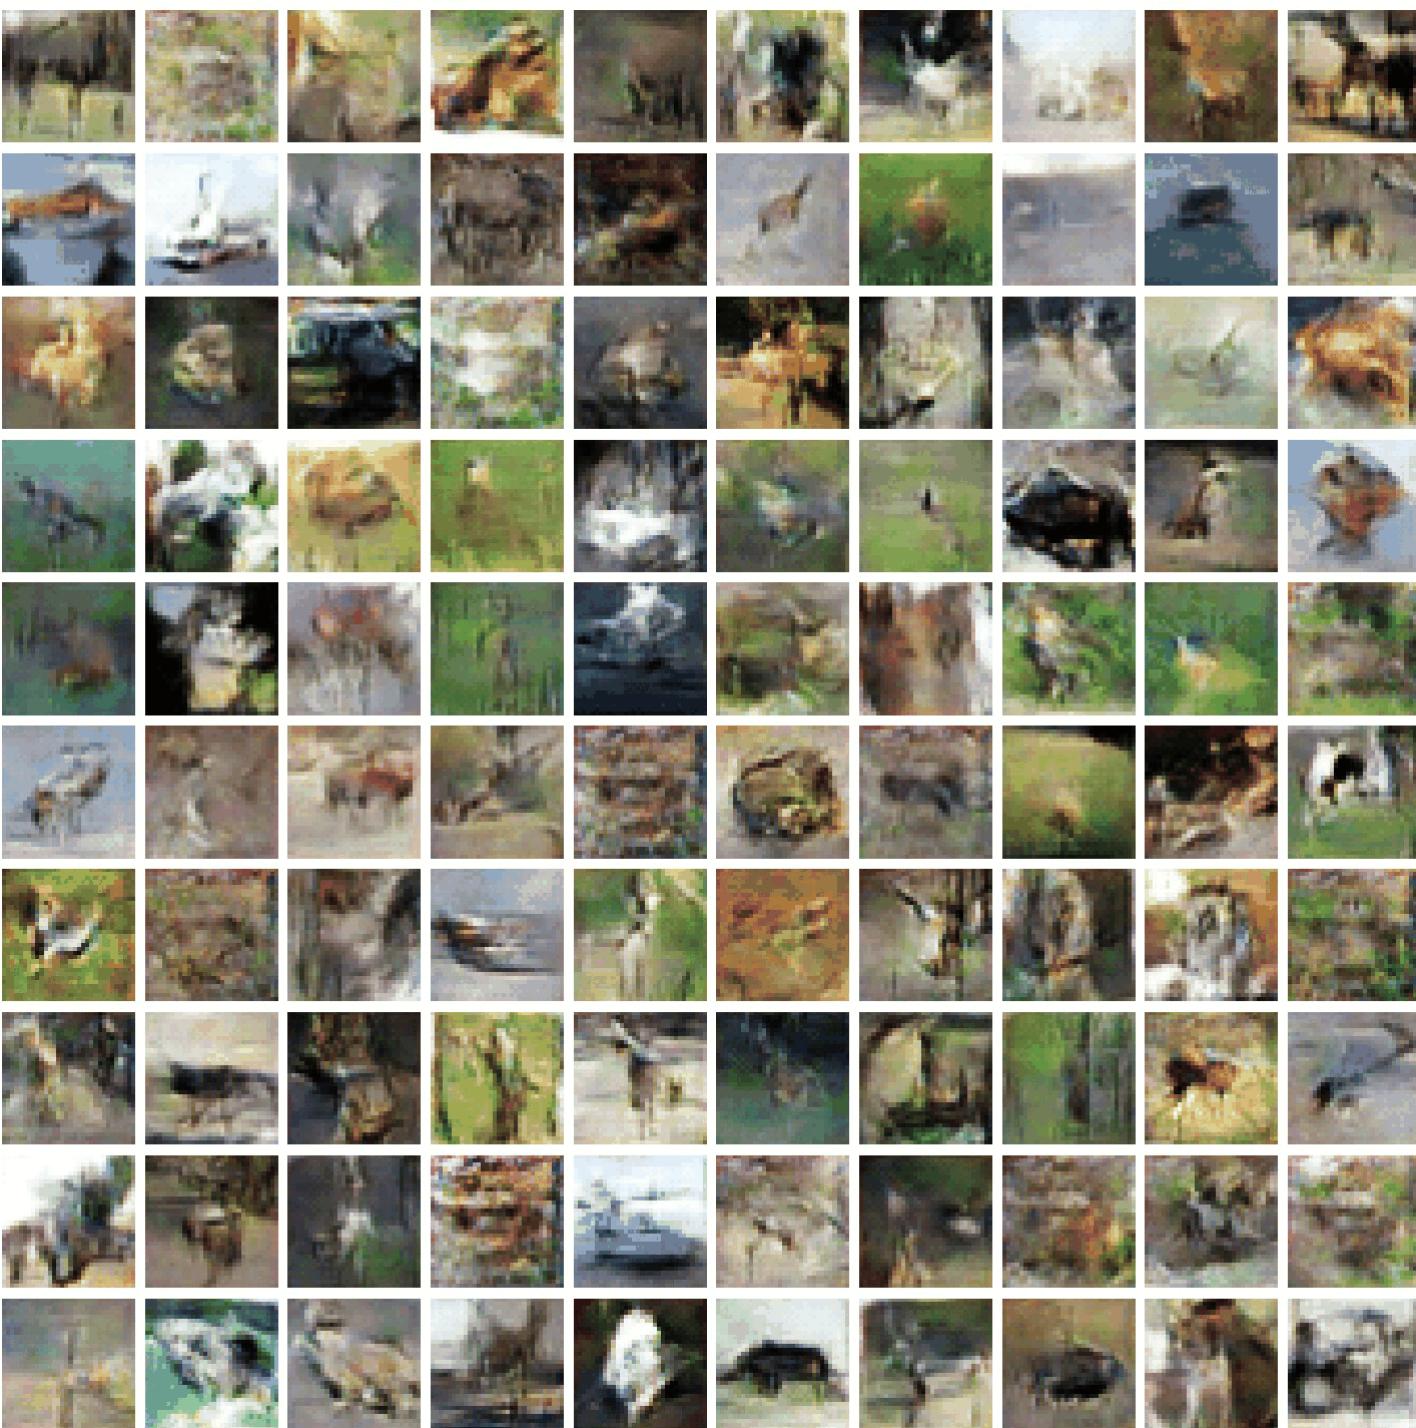
FID Score (Epoch 40): 138.33 ± 0.837



Deep Convolutional GAN Quantitative Results

Architecture	Epoch Checkpoint	FID	KL Divergence	ResNet34 Accuracy
DCGAN	40	138.33 +- 0.837	5.6873	-





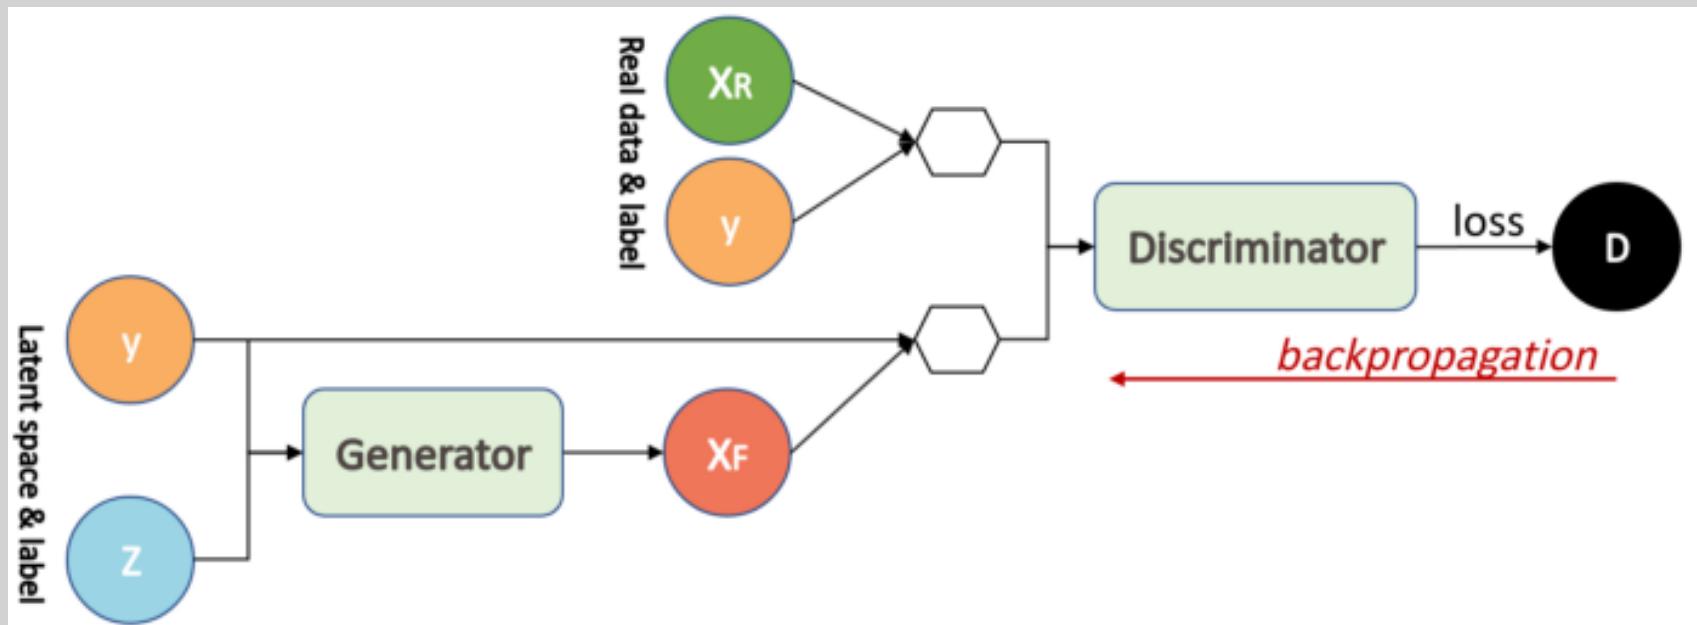
Deep Convolutional GAN Conclusion and Further Improvements

- FID Score (Epoch 40): 138.33 ± 0.837
- FID is quite far from the benchmarks on Papers with Code
- Stable for 40 epochs before encountering a collapse.
- Generated images are not in clear quality, it is hard to identify the focal point of the image.
- Some images were found to be mixed with each other, like a cat and dog was transfigured together into a single image.
- Spectral Normalization and Weights Initialization could have too much regularization effect on the networks.
- Adding Labels into training could improve training and identification. (Conditional GAN, InfoGAN, ACGAN)
- Changing Loss Function could perhaps improve results. (WGAN, WGAN-GP, LSGAN)

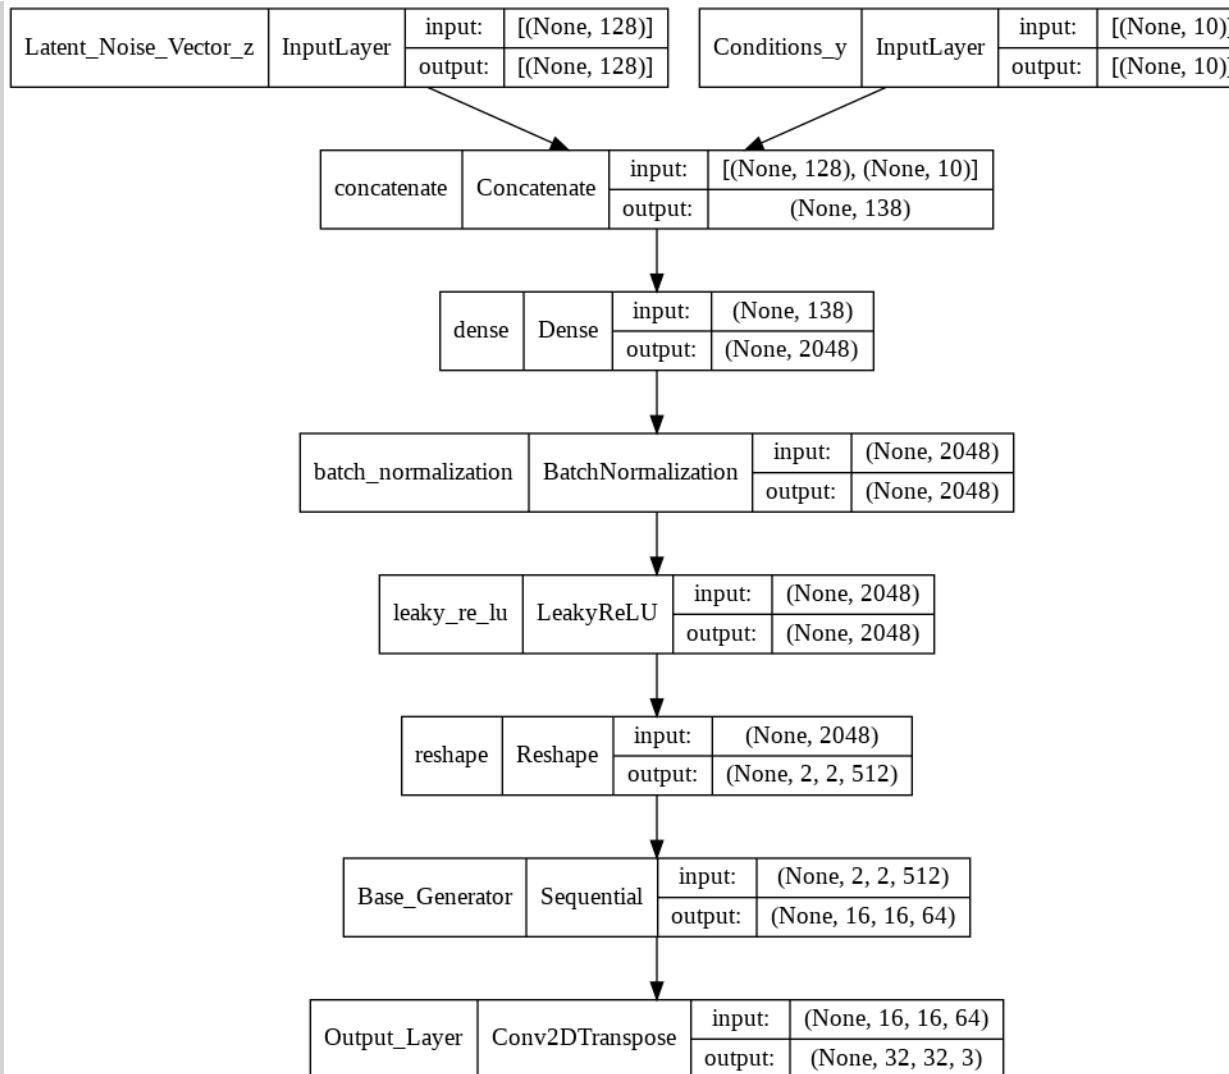


Conditional GAN

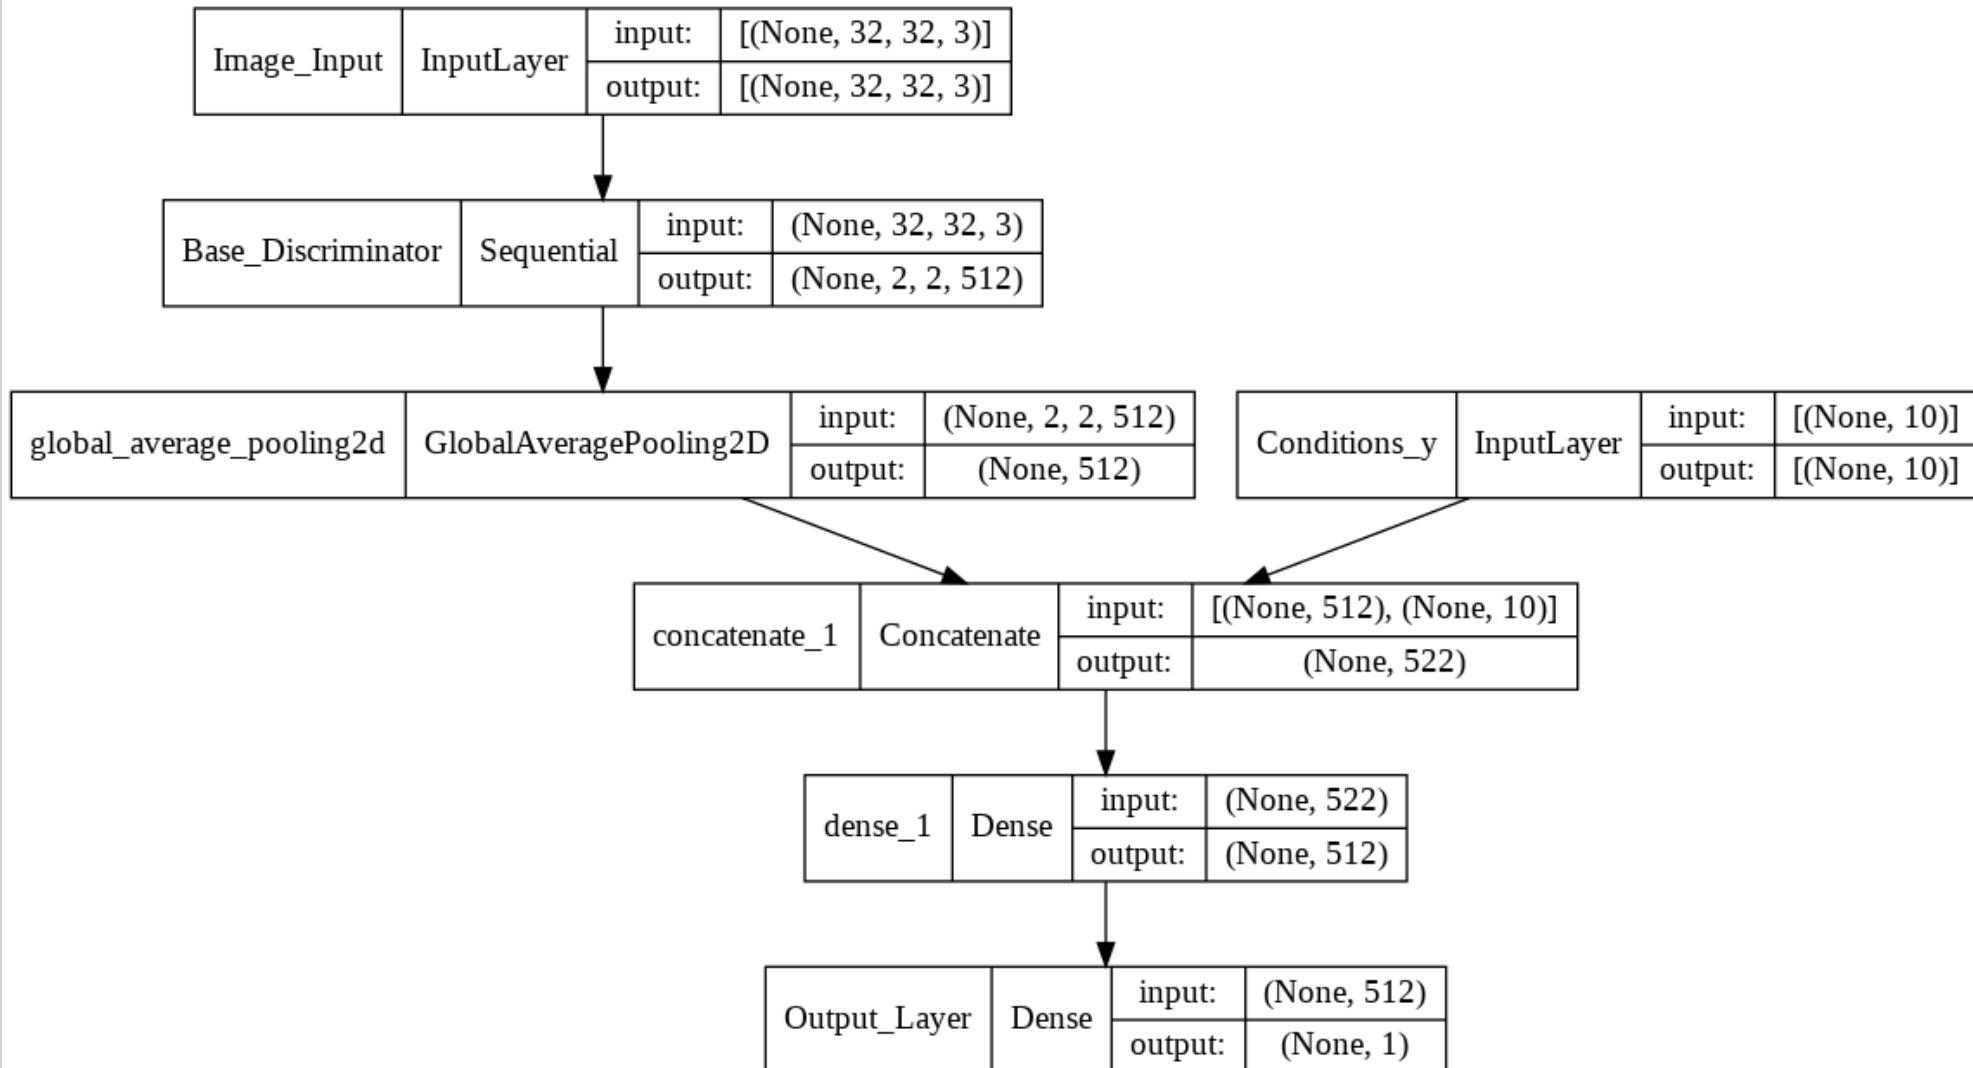
In a Generative Adversarial Network (unconditional), there is no control of the type of images generated. Simply put, it is providing an input to the Generative model and receiving a batch of randomly generated images. These types of GAN are not ideal for the CIFAR-10 dataset, as there are labels provided and we might as well make the full use of it.



Conditional GAN Generator



Conditional GAN Discriminator



Conditional GAN

Epoch 0

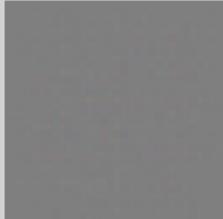
Airplane



Automobile



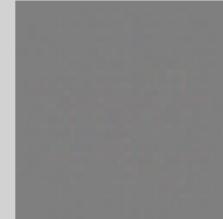
Bird



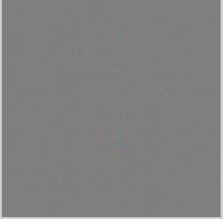
Cat



Deer



Dog



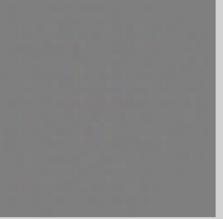
Frog



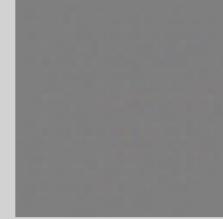
Horse



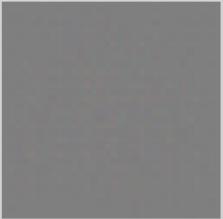
Ship



Truck



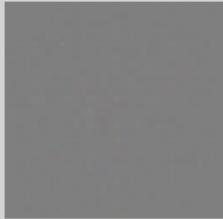
Airplane



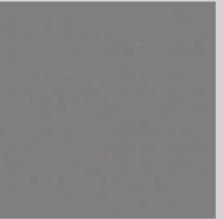
Automobile



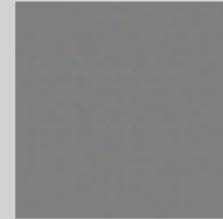
Bird



Cat



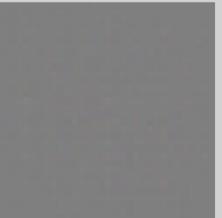
Deer



Dog



Frog



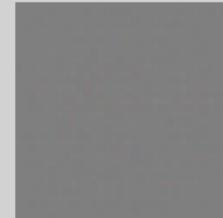
Horse



Ship



Truck



Conditional GAN

Epoch 20

Airplane



Automobile



Bird



Cat



Deer



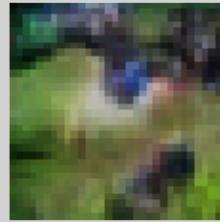
Dog



Frog



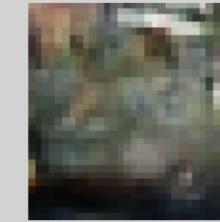
Horse



Ship



Truck



Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Conditional GAN

Epoch 40

Airplane



Automobile



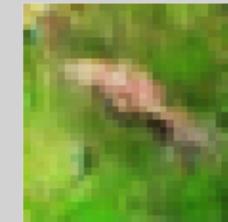
Bird



Cat



Deer



Dog



Frog



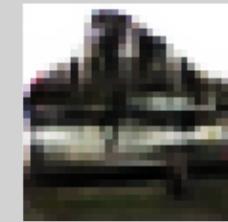
Horse



Ship



Truck



Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Conditional GAN

Epoch 60

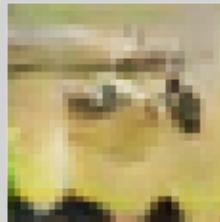
Airplane



Automobile



Bird



Cat



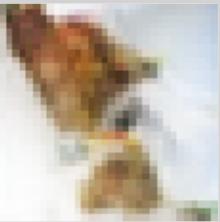
Deer



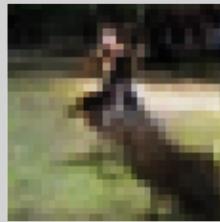
Dog



Frog



Horse



Ship



Truck



Airplane



Automobile



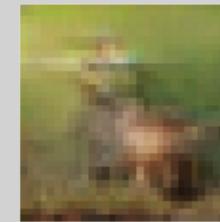
Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Conditional GAN

Epoch 80

Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Conditional GAN

Epoch 100

Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Airplane



Automobile



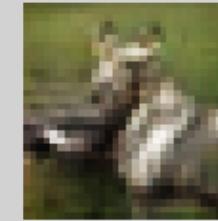
Bird



Cat



Deer



Dog



Frog



Horse



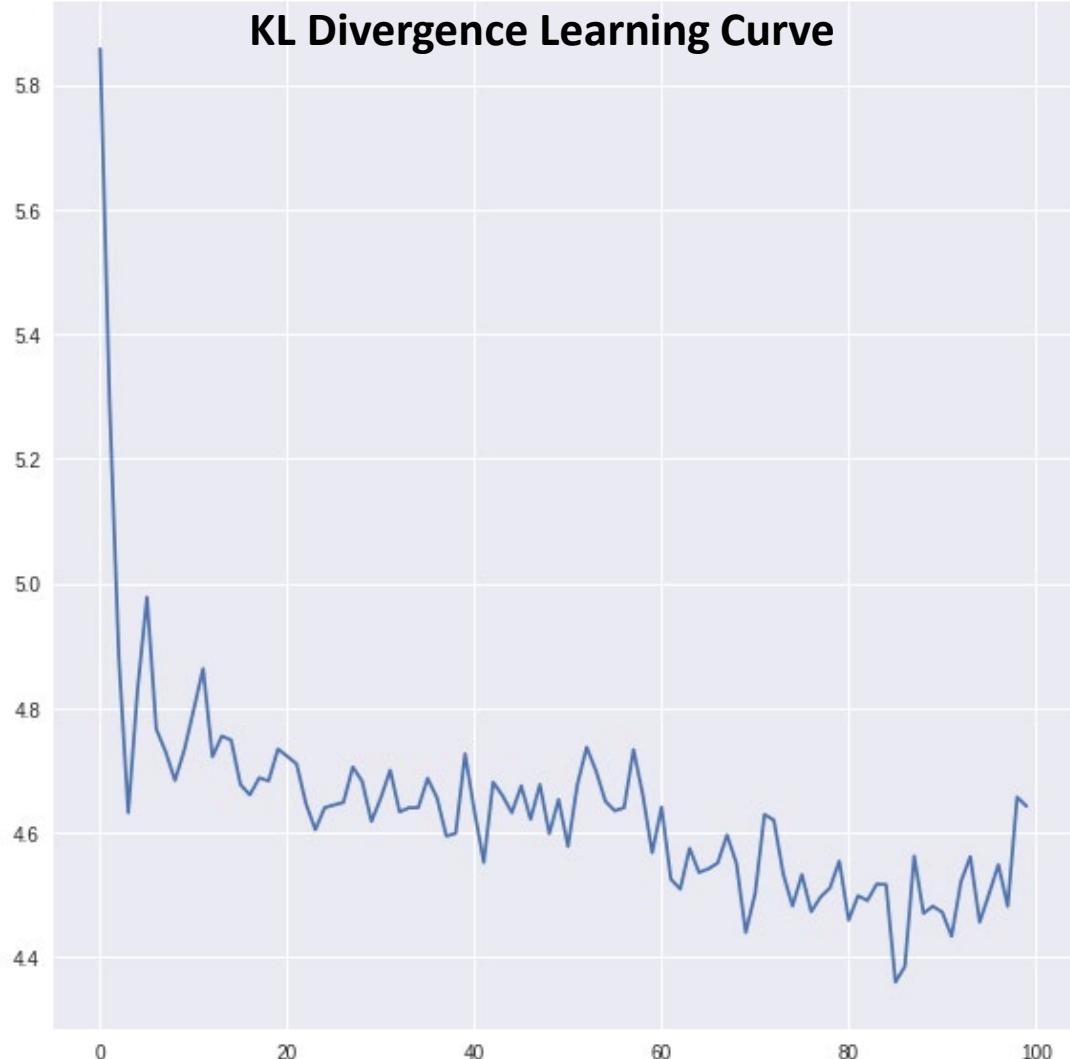
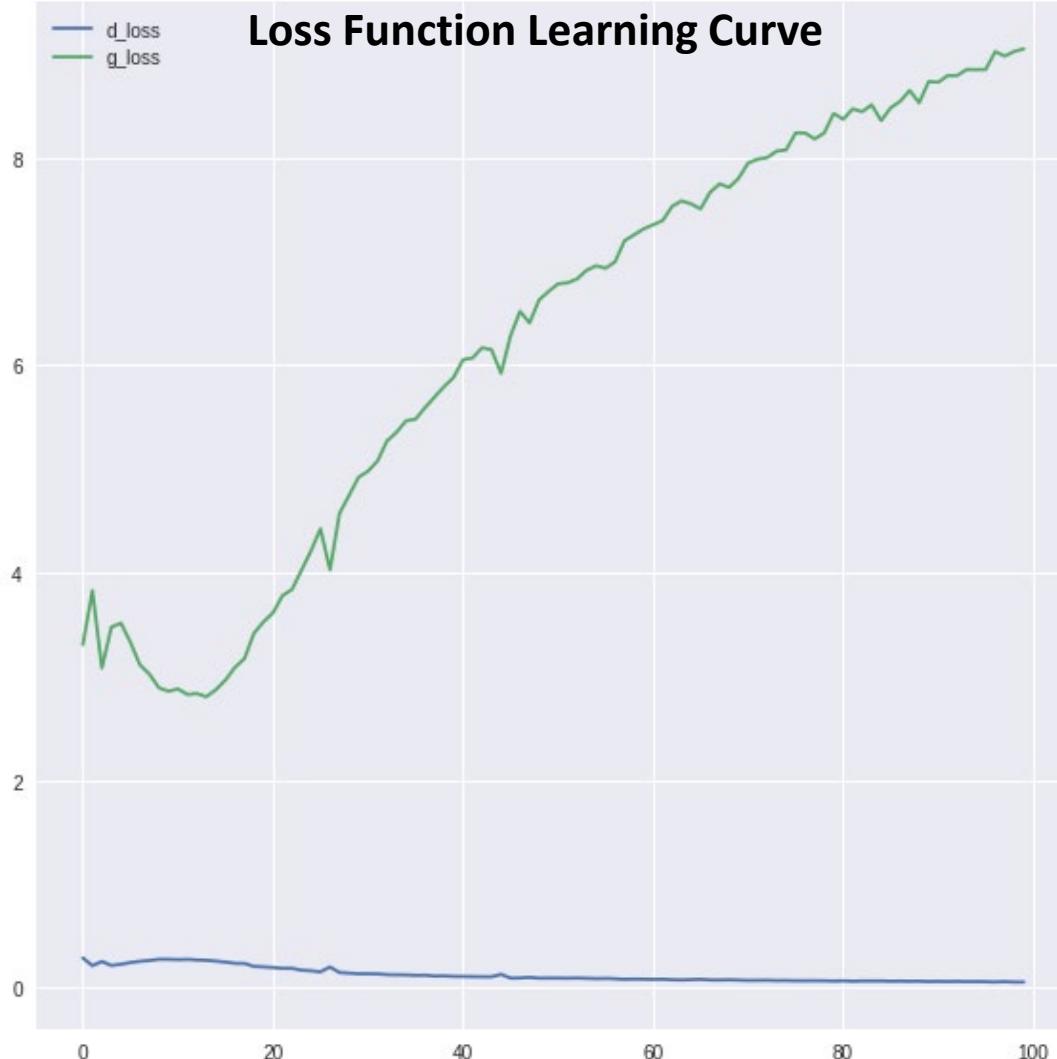
Ship



Truck



Conditional GAN Learning Curve



Conditional GAN

ResNet32 Evaluation

In my opinion, using FID may not be entirely accurate in terms of evaluating our synthetic images, given that the pretrained weights from Inception were trained on the ImageNet dataset. Another issue was that I couldn't tell which class is generated better than the other classes.

In order to resolve these issues, I've pointed out, I will be using an image classifier I've built during ST1501 based on the CIFAR-10 dataset as a surrogate model. The intuition would be using a pre-trained classifier to classify which generated image is like the ones used in the training data. By using this classifier, I can gain a better intuition about the performance of my GAN and which classes are lacking behind.



Conditional GAN

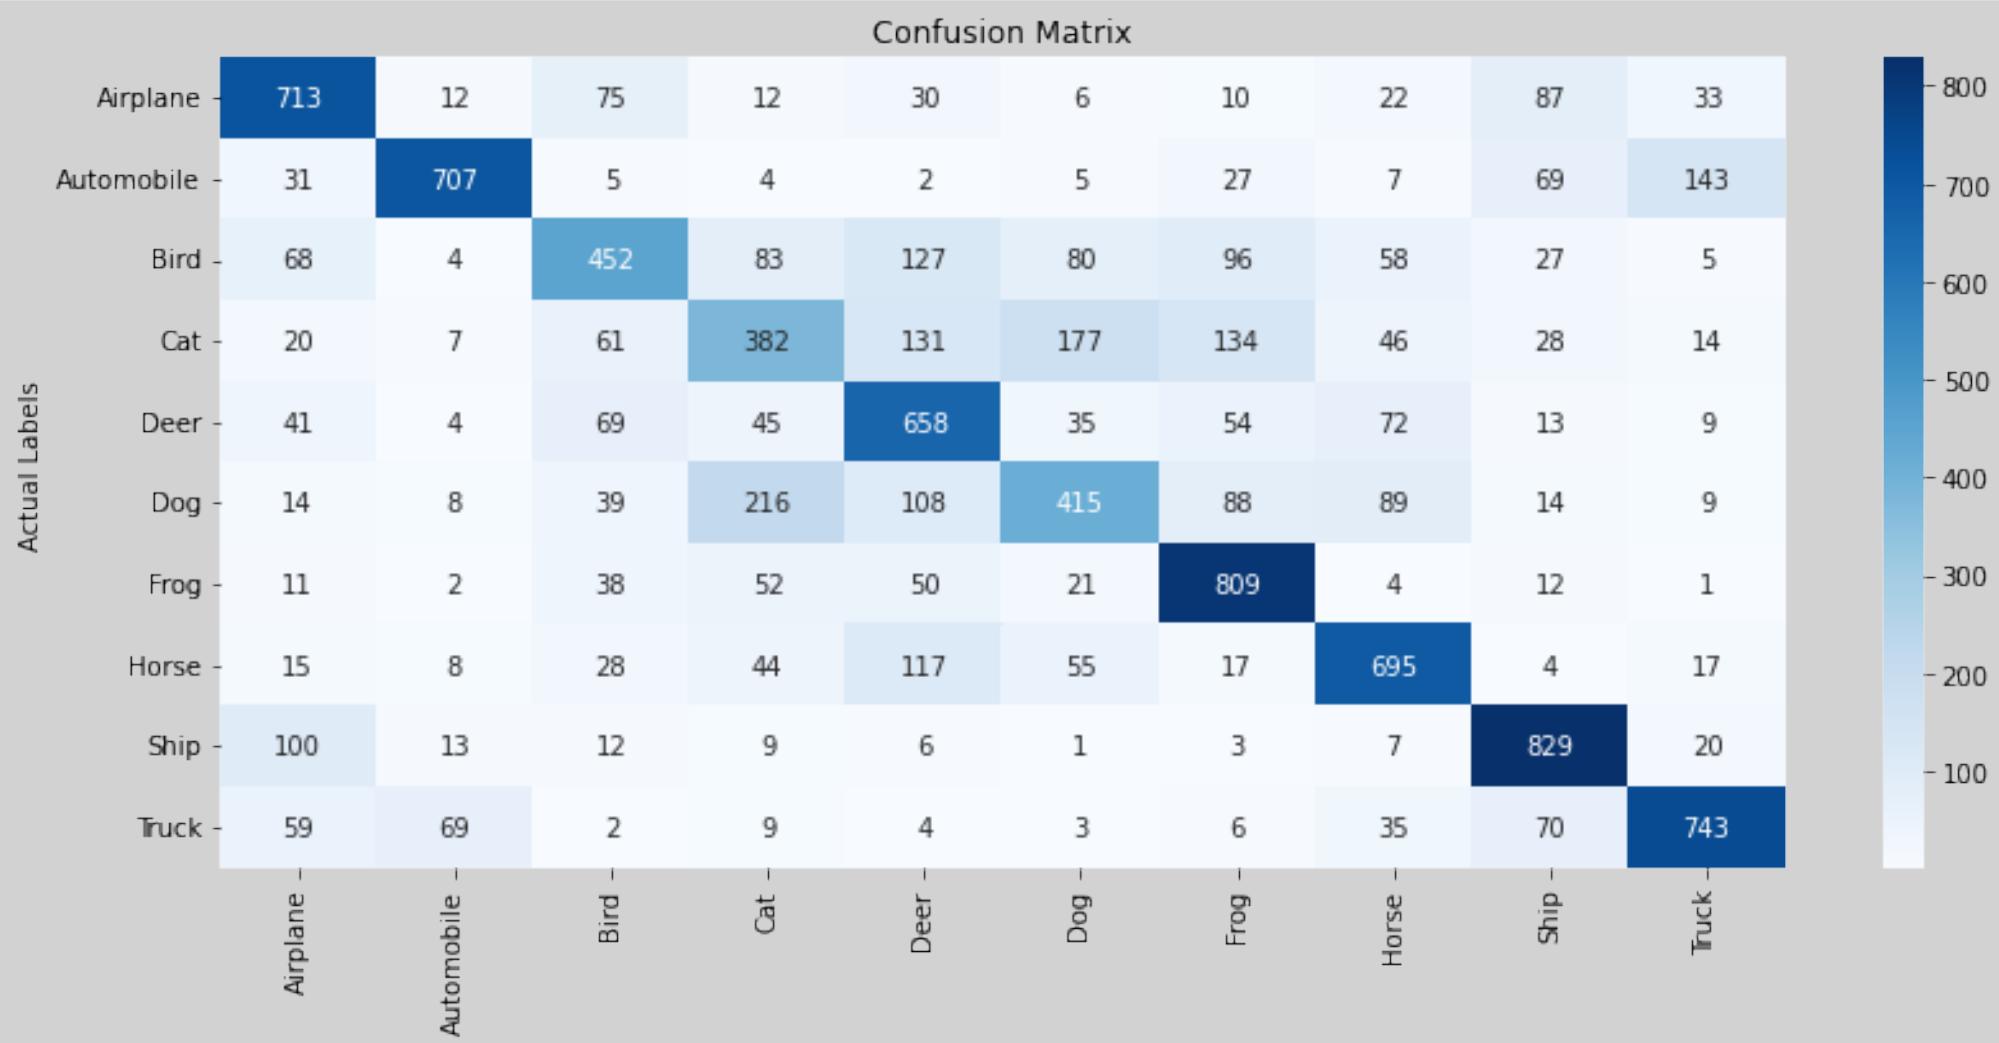
ResNet32 Evaluation

	precision	recall	f1-score	support
Airplane	0.67	0.71	0.69	1000
Automobile	0.85	0.71	0.77	1000
Bird	0.58	0.45	0.51	1000
Cat	0.45	0.38	0.41	1000
Deer	0.53	0.66	0.59	1000
Dog	0.52	0.41	0.46	1000
Frog	0.65	0.81	0.72	1000
Horse	0.67	0.69	0.68	1000
Ship	0.72	0.83	0.77	1000
Truck	0.75	0.74	0.75	1000
accuracy			0.64	10000
macro avg	0.64	0.64	0.63	10000
weighted avg	0.64	0.64	0.63	10000



Conditional GAN

ResNet32 Evaluation



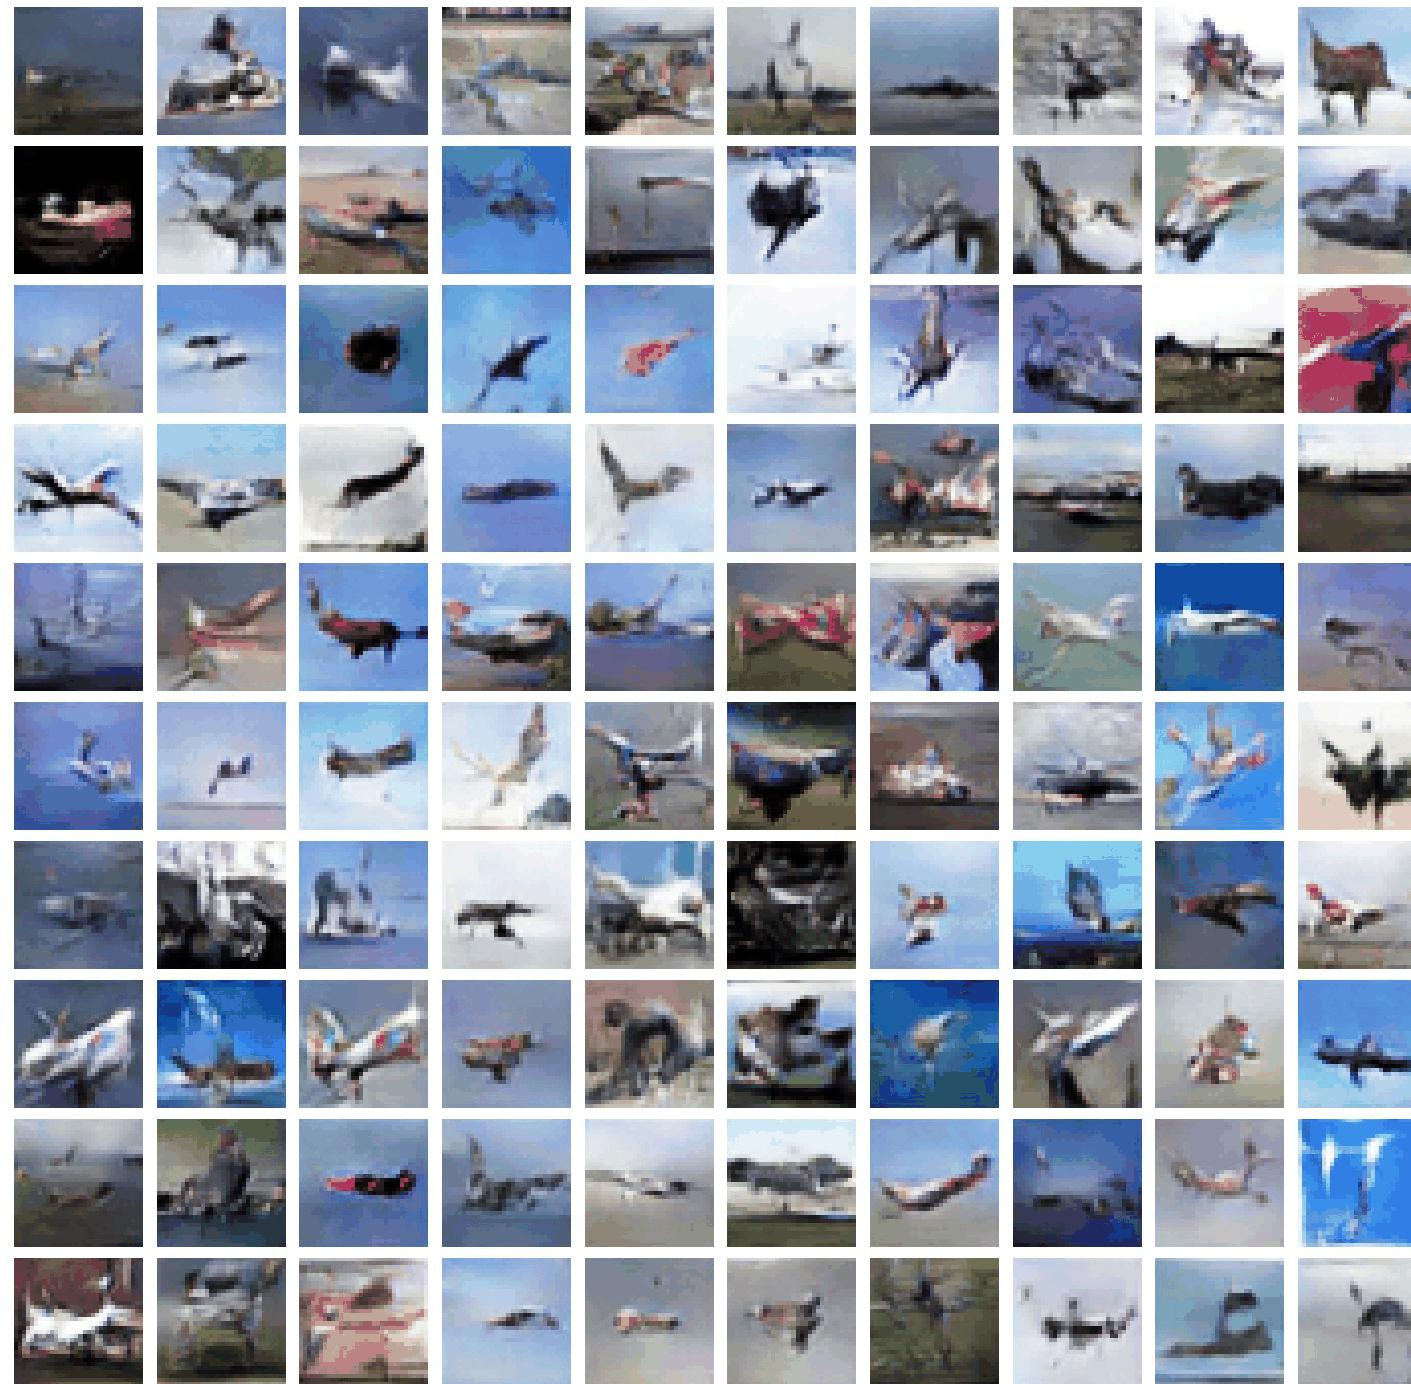
Conditional GAN

Quantitative Results

Architecture	Epoch Checkpoint	FID	KL Divergence	ResNet34 Accuracy
CGAN	100	90.918 +- 0.325	4.6261	0.64



airplane



Conditional GAN Conclusion

- Network was able to train until the pre-set 100 epochs without any issues.
- Fréchet Inception Distance (10,000 sample): 90.918 ± 0.325
- ResNet32 seems to predict only 64% of the synthetic images accurately.
- Majority of **Ship**, **Frogs**, and **Truck** are correctly identified by the image classifier; indicates Generator can generate images of these classes very well.
- Majority of Cat, Dog, and Bird are misclassified by the image classifier; indicates Generator is unable to generate images of these classes very well.
- Animal classes (e.g. Bird, Cat, Deer, Dog) has a tendency to have a false positive error, indicates Generator has difficulty producing images with animal class.



Conditional GAN

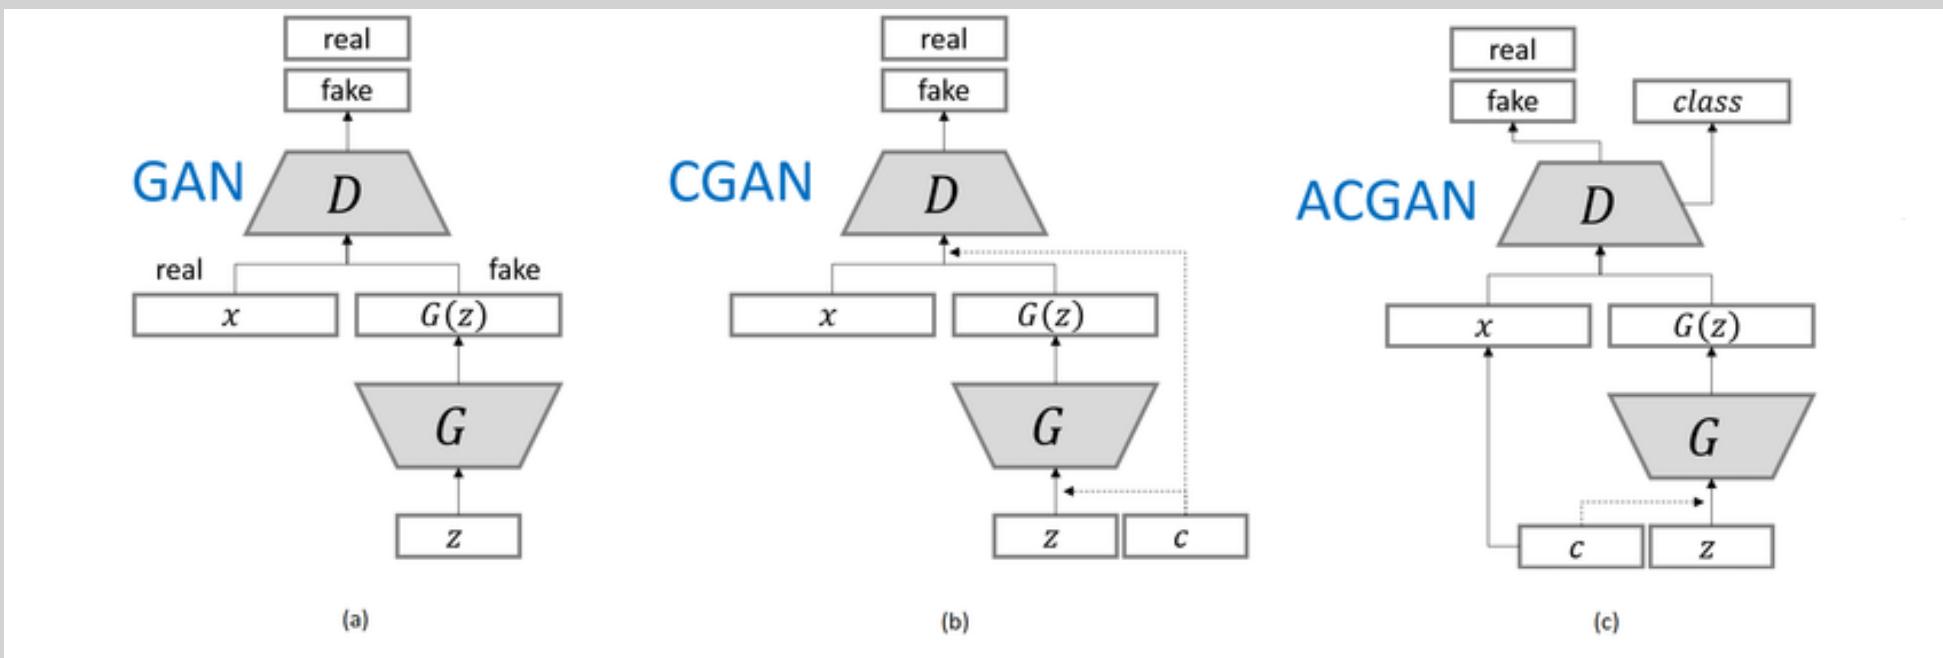
Further Improvements

- Perhaps making better use of the Labels could improve image generation quality (ACGAN, Semi-Supervised GAN)
- Changing Loss Functions, e.g. Hinge Loss, Wasserstein Loss
- Label Smoothing
- Multiple Generator Update per Discriminator Update
- Differentiable Augmentation

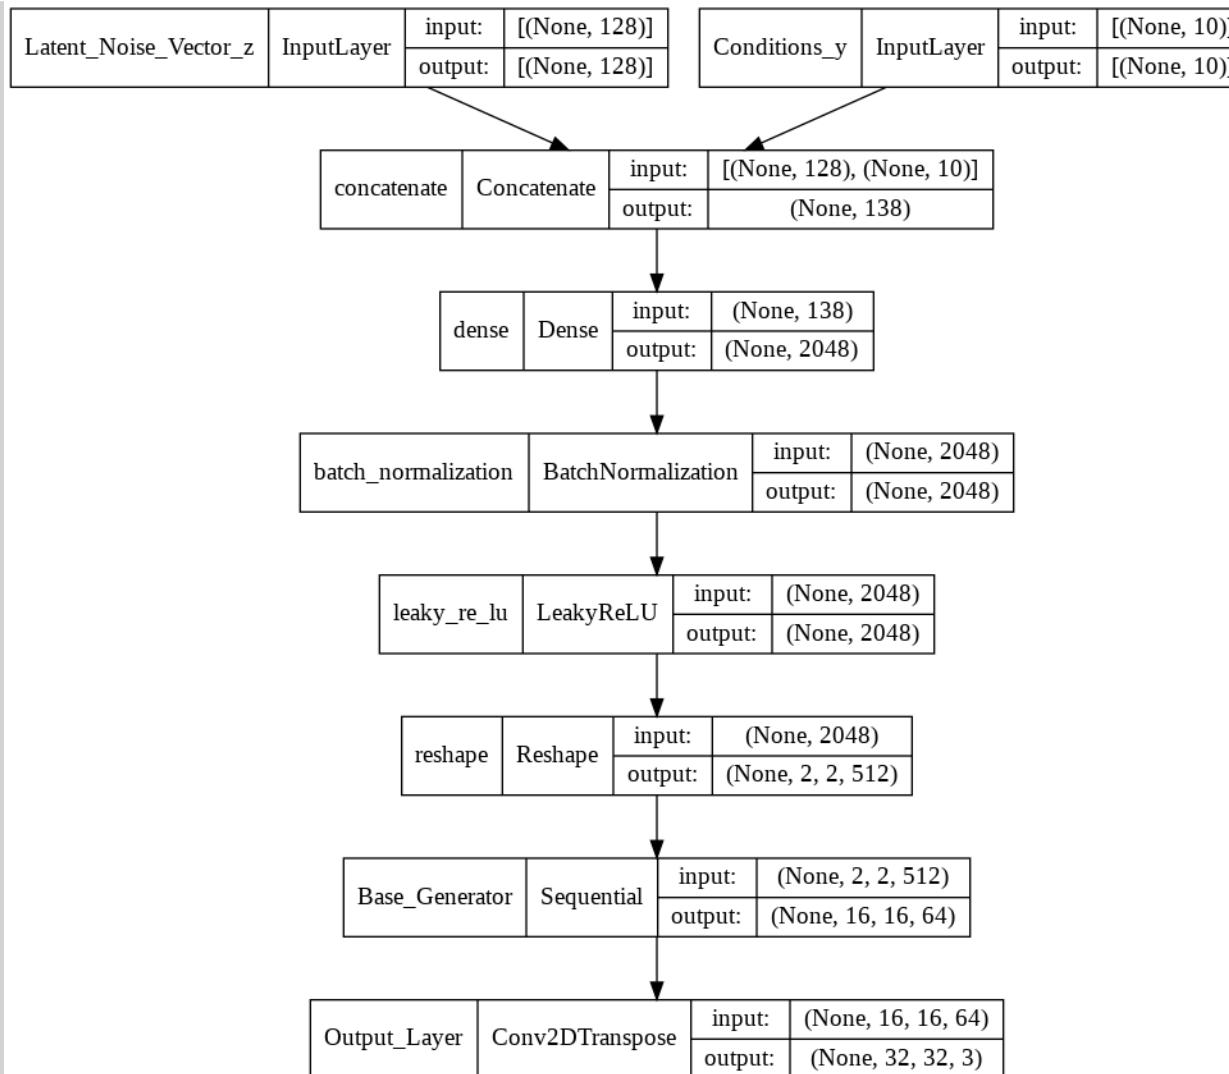


Auxiliary Classifier GAN

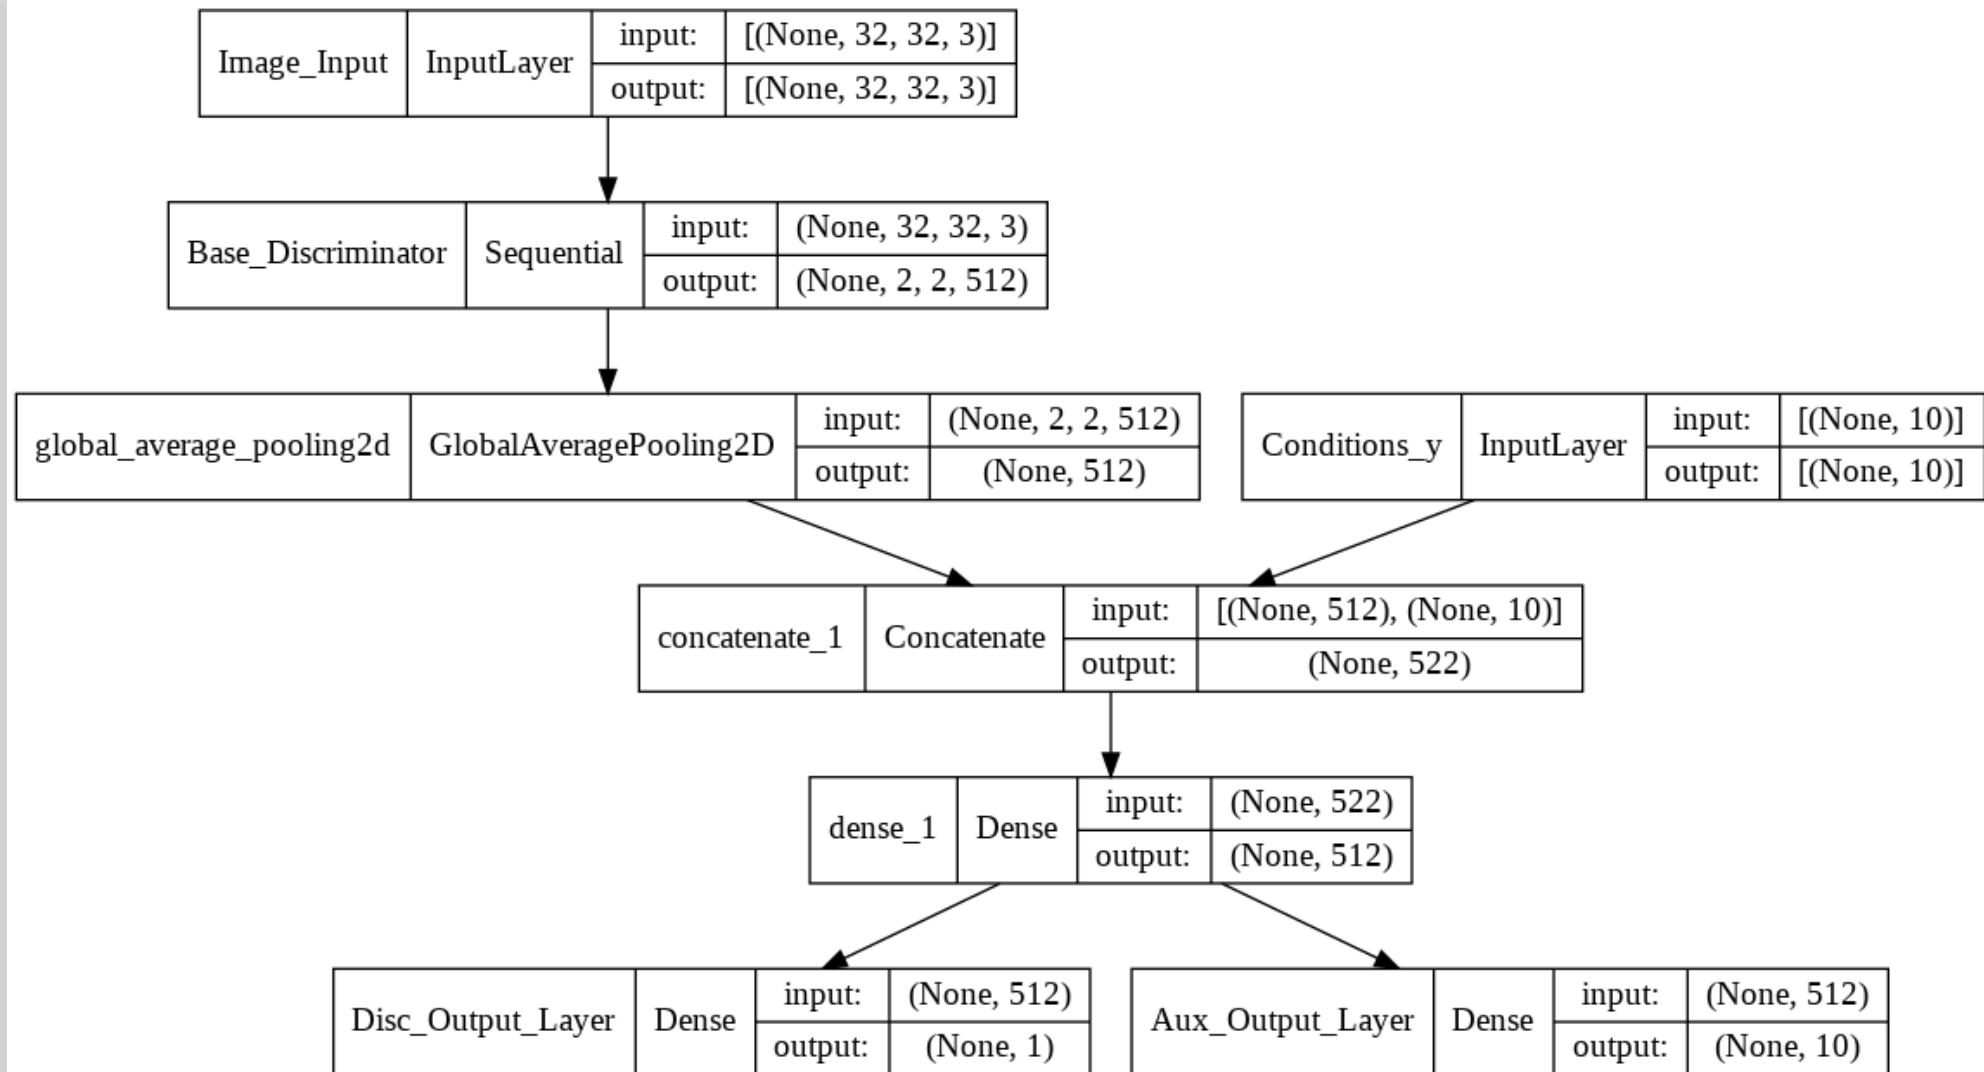
Auxiliary Classifier Generative Adversarial Network, or ACGAN in short, follows a similar architecture to a Conditional GAN. In fact, I've only added less than 10 lines of code for this implementation. There are two main difference between the two architectures - the discriminator produces two outputs, and an auxiliary loss function is included in the training process.



Auxiliary Classifier GAN Generator



Auxiliary Classifier GAN Discriminator



Auxiliary Classifier GAN Loss Function

There is an important concept for AC-GAN during training process. Instead of having one loss function for the discriminator, there are two loss function for different purpose.

- **Discriminator Loss:** Error between actual image and generated image (Binary Cross Entropy)
- **Auxiliary Classifier Loss:** Error between actual conditions and predicted conditions (Categorical Cross Entropy)

Total Loss = Discriminator Loss + Auxiliary Classifier Loss





Auxiliary Classifier GAN Training Function

```
● ● ●

# Train the discriminator.
with tf.GradientTape() as tape:
    disc_output, aux_output = self.discriminator([combined_images, combined_condition])
    disc_d_loss = self.disc_loss_fn(labels, disc_output)
    aux_d_loss = self.aux_loss_fn(combined_condition, aux_output)
    d_loss = disc_d_loss + aux_d_loss
    grads = tape.gradient(d_loss, self.discriminator.trainable_weights)
    self.d_optimizer.apply_gradients(
        zip(grads, self.discriminator.trainable_weights)
    )

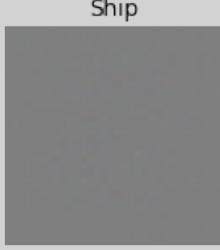
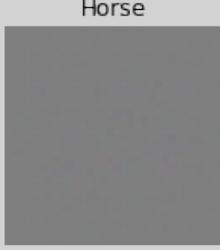
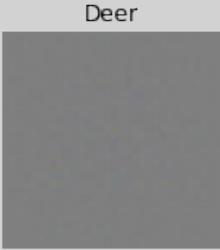
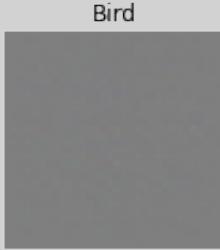
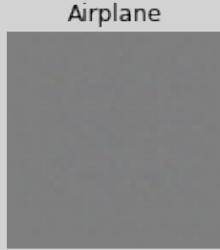
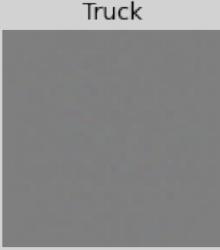
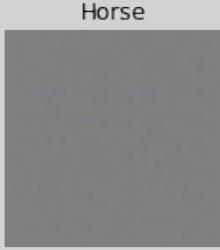
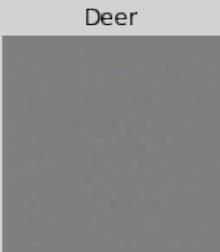
● ● ●

# Train the Generator
with tf.GradientTape() as tape:
    fake_images = self.generator([latent_noise_vector, condition])
    disc_output, aux_output = self.discriminator([fake_images, condition])
    disc_g_loss = self.disc_loss_fn(misleading_labels, disc_output)
    aux_g_loss = self.aux_loss_fn(condition, aux_output)
    g_loss = disc_g_loss + aux_g_loss
    grads = tape.gradient(g_loss, self.generator.trainable_weights)
    self.g_optimizer.apply_gradients(zip(grads, self.generator.trainable_weights))
```



Auxiliary Classifier GAN

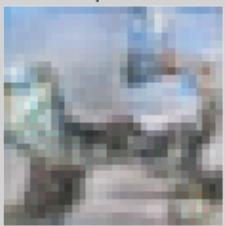
Epoch 0



Auxiliary Classifier GAN

Epoch 20

Airplane



Automobile



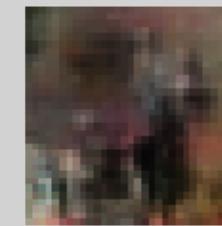
Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



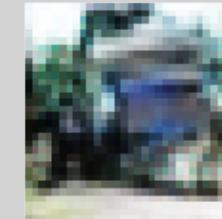
Horse



Ship



Truck



Auxiliary Classifier GAN

Epoch 40

Airplane



Automobile



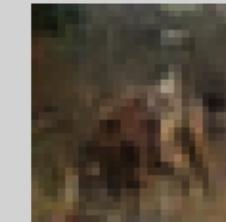
Bird



Cat



Deer



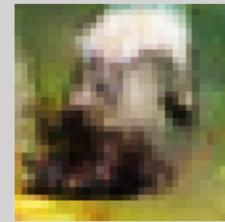
Dog



Frog



Horse



Ship



Truck



Airplane



Automobile



Bird



Cat



Deer



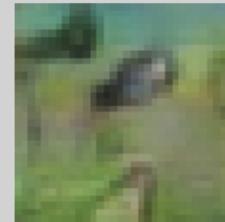
Dog



Frog



Horse



Ship



Truck



Auxiliary Classifier GAN

Epoch 60

Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



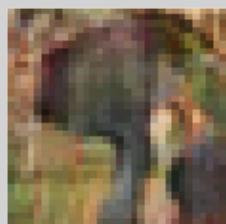
Airplane



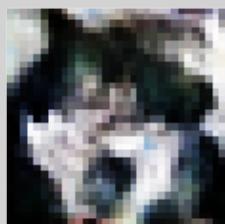
Automobile



Bird



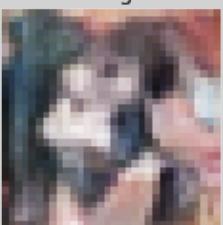
Cat



Deer



Dog



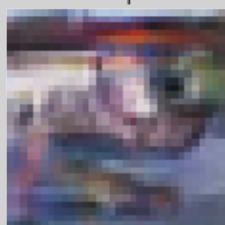
Frog



Horse



Ship



Truck



Auxiliary Classifier GAN

Epoch 80

Airplane



Automobile



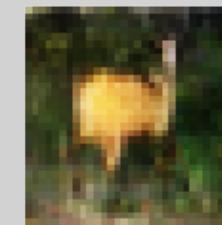
Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck



Auxiliary Classifier GAN

Epoch 100

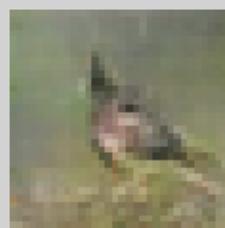
Airplane



Automobile



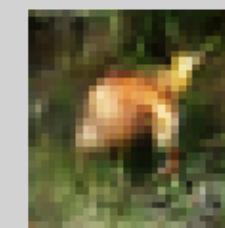
Bird



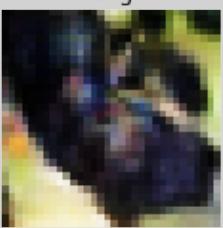
Cat



Deer



Dog



Frog



Horse



Ship



Truck



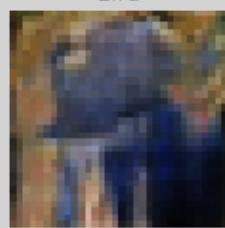
Airplane



Automobile



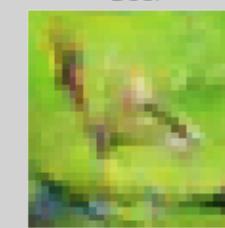
Bird



Cat



Deer



Dog



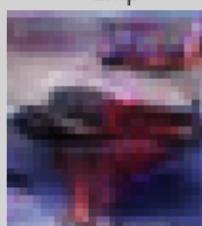
Frog



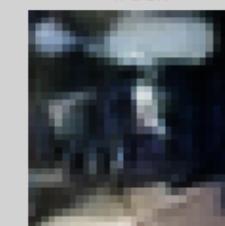
Horse



Ship

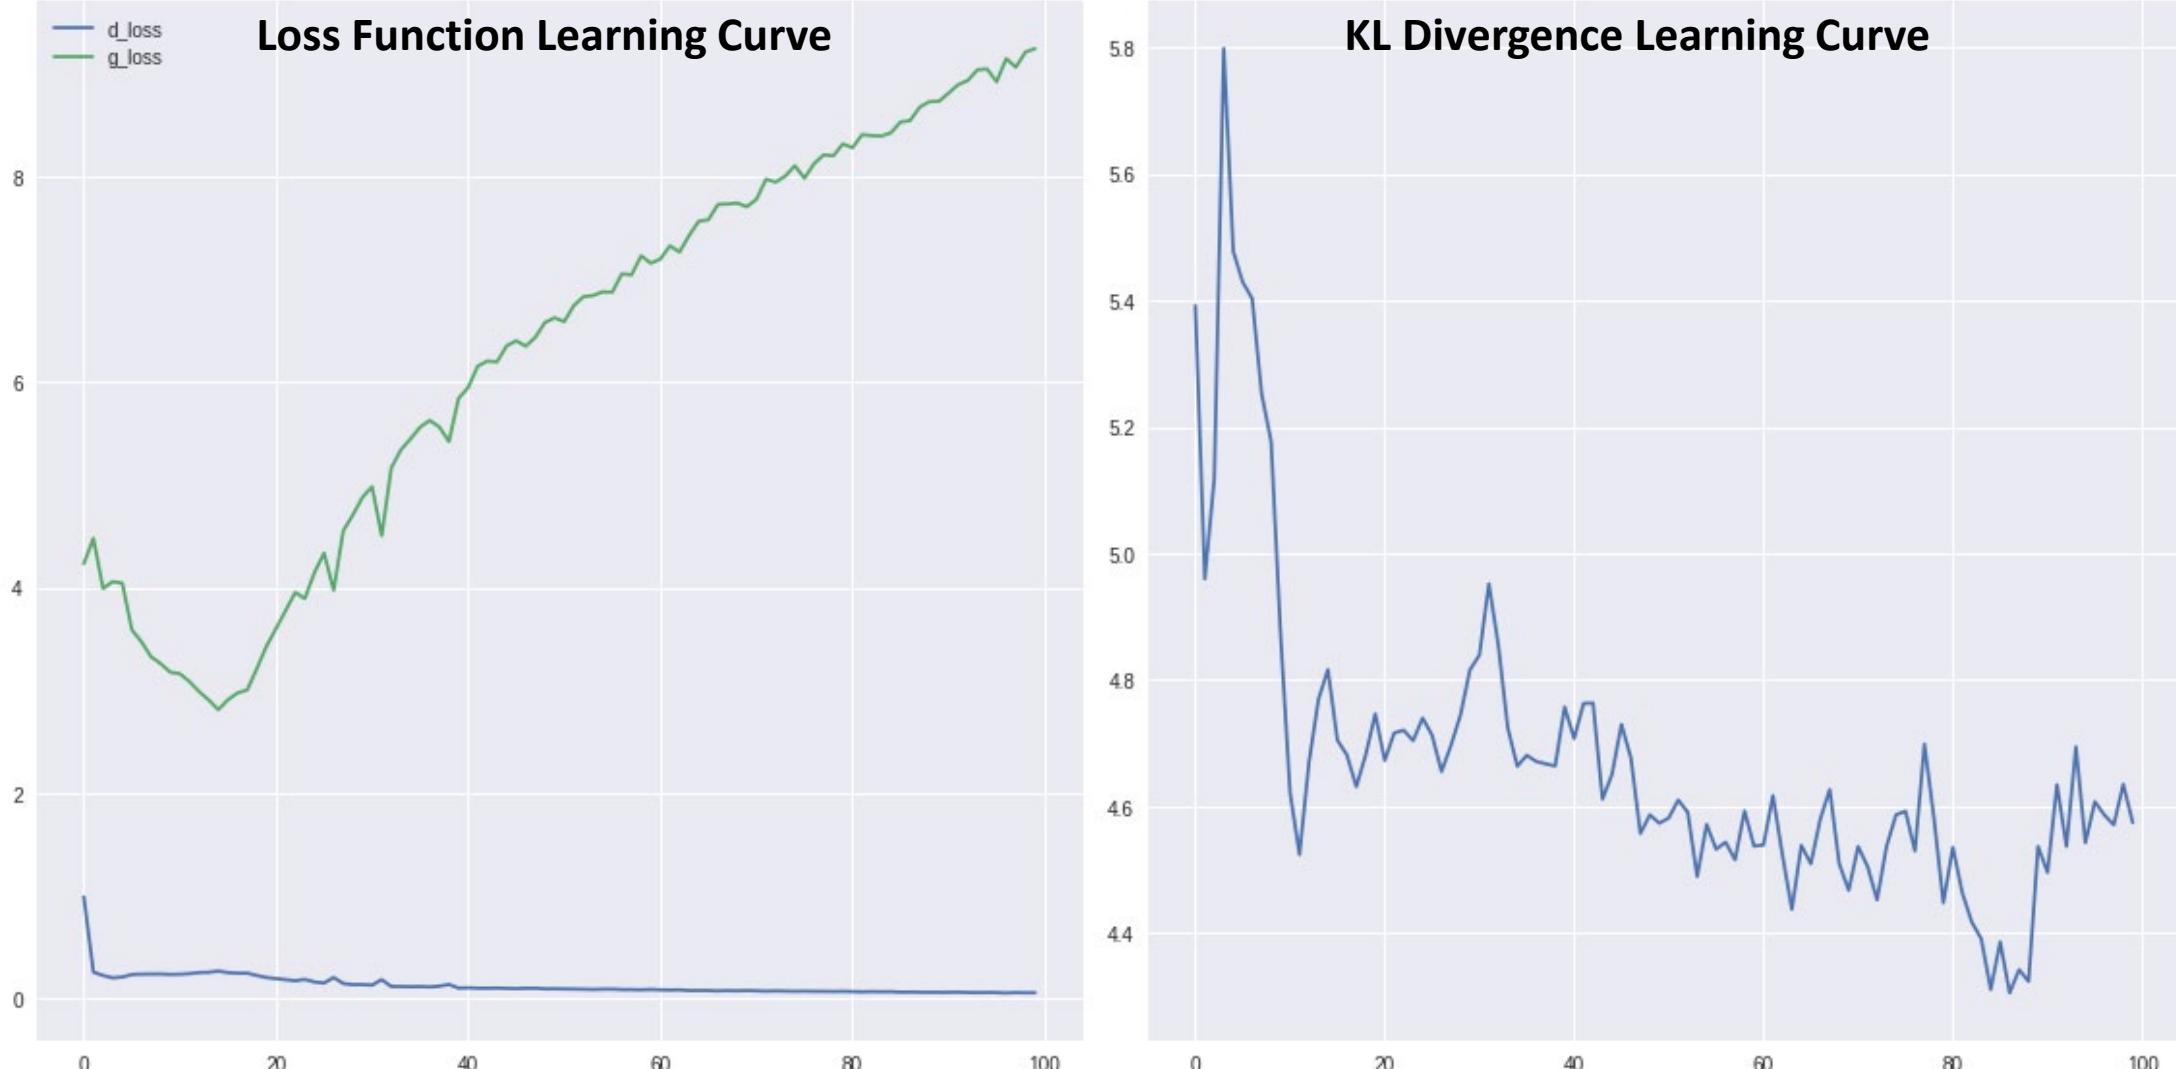


Truck



Auxiliary Classifier GAN

Learning Curve



Auxiliary Classifier GAN

ResNet32 Evaluation

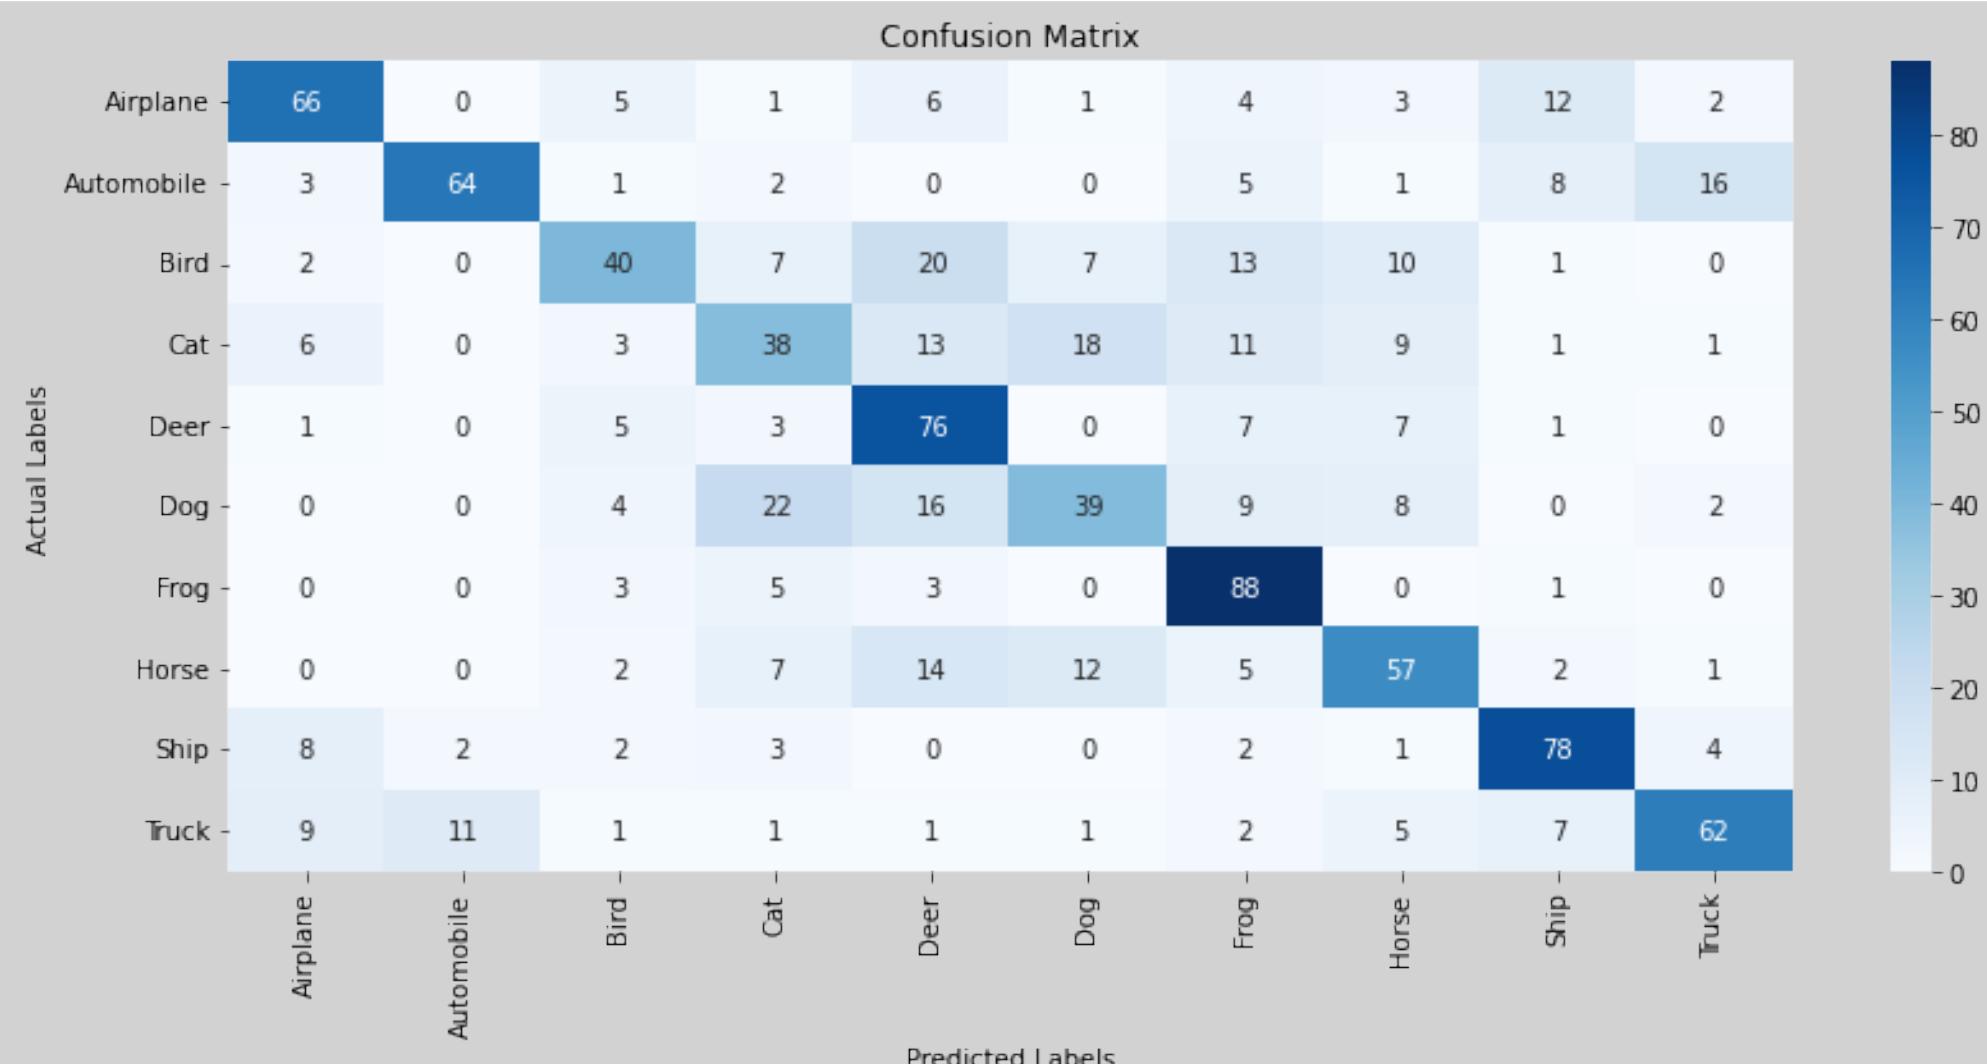
	precision	recall	f1-score	support
Airplane	0.69	0.66	0.68	100
Automobile	0.83	0.64	0.72	100
Bird	0.61	0.40	0.48	100
Cat	0.43	0.38	0.40	100
Deer	0.51	0.76	0.61	100
Dog	0.50	0.39	0.44	100
Frog	0.60	0.88	0.72	100
Horse	0.56	0.57	0.57	100
Ship	0.70	0.78	0.74	100
Truck	0.70	0.62	0.66	100
accuracy			0.61	1000
macro avg	0.61	0.61	0.60	1000
weighted avg	0.61	0.61	0.60	1000





Auxiliary Classifier GAN

ResNet32 Evaluation



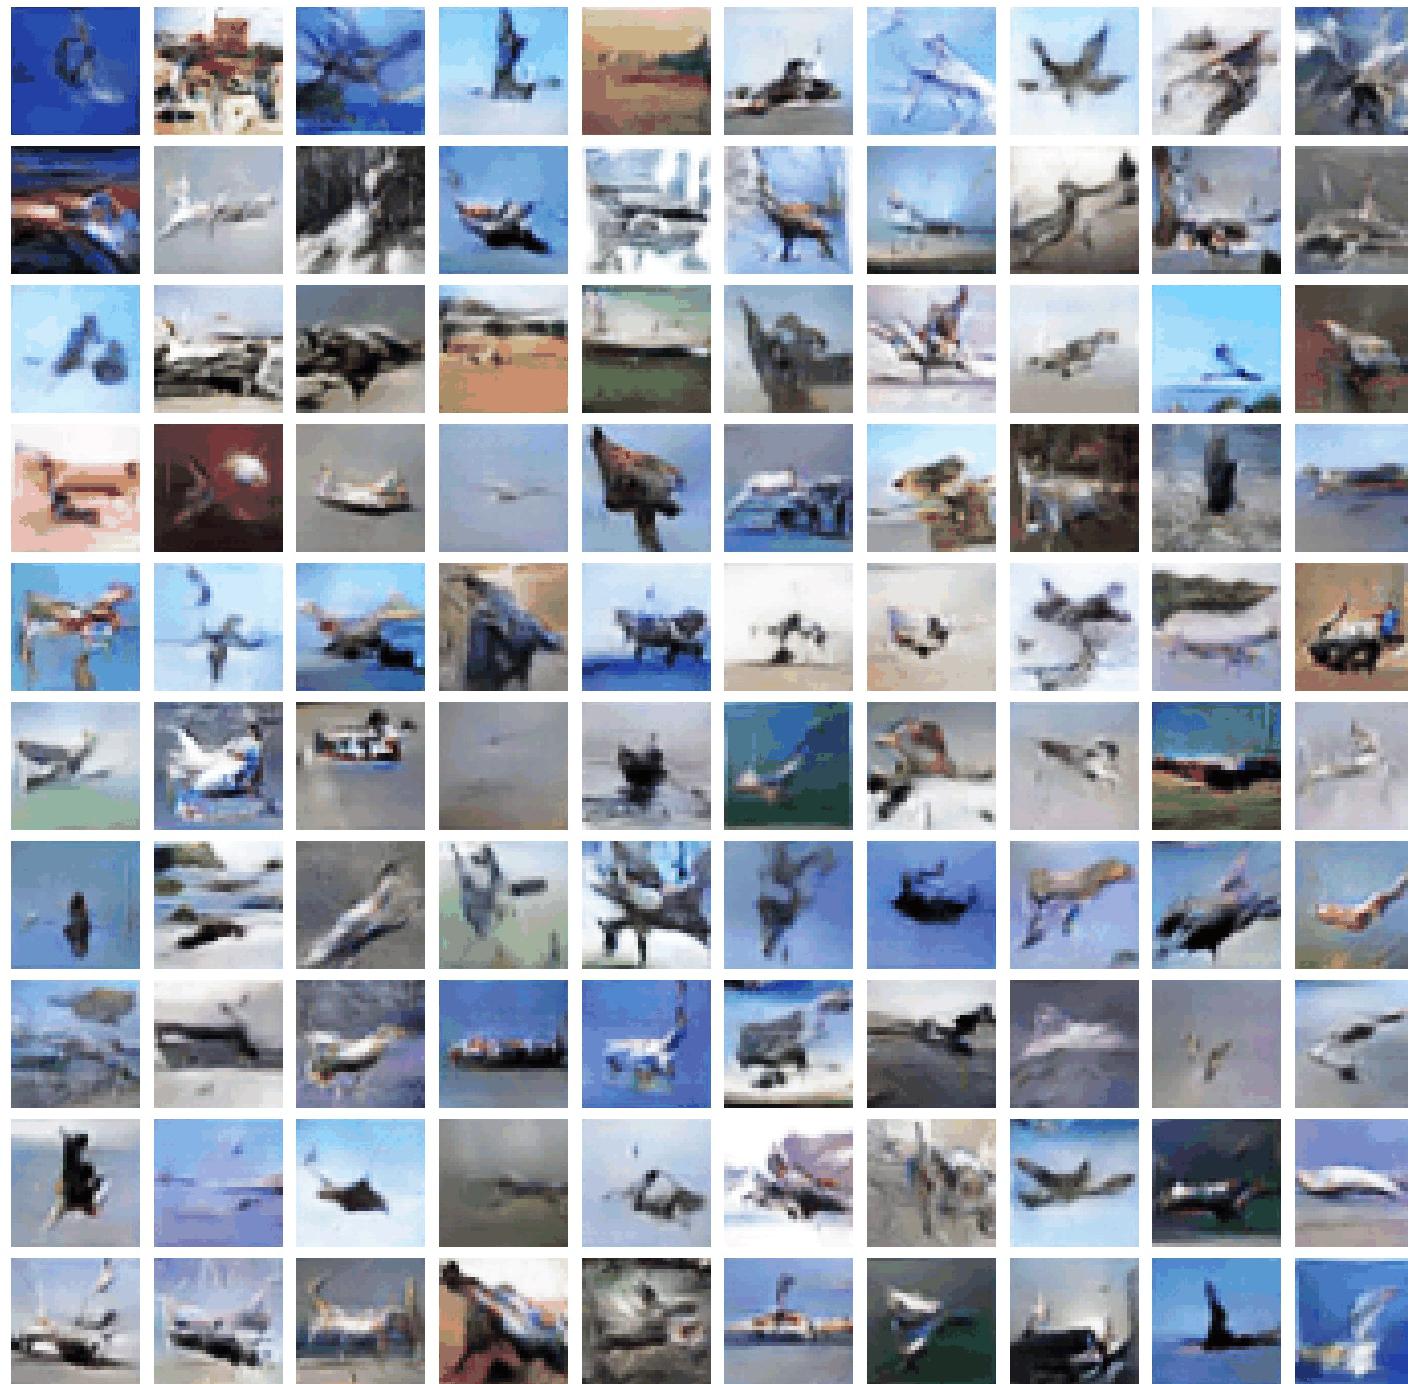
Auxiliary Classifier GAN

Quantitative Results

Architecture	Epoch Checkpoint	FID	KL Divergence	ResNet34 Accuracy
ACGAN	80	93.149 +- 0.737	4.4176	0.61



airplane



Auxiliary Classifier GAN Conclusion

- KL Divergence indicates that the model trains very well until epoch 80 - fail collapse problem.
- Fréchet Inception Distance (10,000 samples): 93.149 ± 0.737
- ResNet32 seems to predict only 61% of the synthetic images accurately.
- Majority of **Frogs**, **Ship**, and **Truck** are correctly identified by the image classifier; indicates Generator can generate images of these classes very well.
- Majority of **Bird**, **Dog**, and **Cat** are misclassified by the image classifier; indicates Generator is unable to generate images of these classes very well.
- Animal classes (e.g. Bird, Cat, Deer, Dog) tends to have a false positive error, indicates Generator has difficulty producing images with animal class.



Auxiliary Classifier GAN

Further Improvement

- Attempt to change loss function, e.g. Hinge Loss, Wasserstein Loss
- Label Smoothing
- Multiple Generator Update per Discriminator Update
- Differentiable Augmentation



Final Quantitative Results

Architecture	Epoch Checkpoint	FID	KL Divergence	ResNet34 Accuracy
GAN	40	138.33 +- 0.837	5.6873	-
CGAN	100	90.918 +- 0.325	4.6261	0.64
ACGAN	80	93.149 +- 0.737	4.4176	0.61 

Conclusion

- I have attempted three different GAN architectures – DCGAN, CGAN, ACGAN
- Based on these implementation, I have learnt that every little details in the architecture does matter to training.
- Probably is the reason why there is so many research conducted on GAN.
- Further Improvements that could not have been implemented due to time constraint – changing loss function, label smoothing, multiple time update rule, differentiable augmentation, skip connections, and many more.

